What produces a capturable complexity gap: one representative economy per model and group

Alexander Ober, with Claude

October 11, 2026; restructured October 15, final numbers October 24, loading-size, compensation and non-monotone ladders October 30, 2026

Scope. Every economy with a protocol run (the committed record in `variants/results/economy_table.csv` plus the October 6–15, 2026 Quest campaign in `bop_runs/results`) and the reduced-form screens. The main tables hold one representative per model per group, built so that the same change is made in all three models; everything else that was run is in the appendix (Section 6). The tables are generated by `variants/campaign_tables.py` from the oracle and estimator files. Every campaign economy ran ten seeds (`gvypk3` twenty; the reduced form five per cell).

Reading the tables. Every table has the same columns. $\mathrm{SR}_{\max}$, the *linear ceiling* and the *room* are the population investor's conditional Sharpe ratios over the evaluation window: the linear ceiling is the best Sharpe a model linear in the ranks of the characteristics reaches with the true conditional moments, the room is the best nonlinear ceiling (random features, bins, quadratic) minus that linear ceiling, and *room % of linear* is the room as a percentage of the linear ceiling. That is the theoretical scope for nonlinearity, before any estimation. Sharpe ratios add in quadrature, so the room understates an interaction that sits on a large linear base; *interaction SR* is the nonlinear ceiling's Sharpe orthogonal to the linear one, $\sqrt{\mathrm{nl}^2 - \mathrm{lin}^2}$, the size of the interaction itself, and *DKKM* $\perp$ is DKKM's Sharpe orthogonal to the best classical method, signed like the gap. *Market share* is the equal-weighted market's share of $\mathrm{SR}_{\max}$. *DKKM*, *FMR* (Fama–MacBeth), *FF* (Fama–French factors), *linear* (the best of rank and level regressions, with and without the market) and *CAPM* (the estimated market line) are each method's out-of-sample Sharpe at the 360-month window. All Sharpe ratios are monthly (multiply by $\sqrt{12}$ for annual), so a gap of 0.01 is 0.035 a year. The *fair gap* is DKKM minus the best of the four classical methods, averaged across seeds, with in parentheses its t-statistic (the mean over its standard error across seeds; ten seeds unless stated); "+" counts the seeds in which it is positive. A dagger marks an oracle that predates the evaluation-window ceilings.

1 The answer

Each model has a spine: its published baseline, and one economy in which a priced state y enters firm cash flows through an interaction of two firm variables (group 3b). Every other group is one change to that spine, the same change in all three models. Only the interaction produces a fair gap. With one type, every firm's exposure varies with one firm variable (3a), the premium is monotone in one characteristic's rank, and room is 5 to 10% of the linear ceiling with no gap; a crash-only factor (1) gives room of 3 to 12% and no gap; a quadratic price of the model's own risk (2) gives room of 0 to 14% and no gap; the interaction economies (3b) have room of 16 to 47% and gaps of +0.008 to +0.020 (t 2.7 to 5.2); with the type loading's cross-sectional mean set to zero (`vyg25s`, `bvychs`, Section 6) room rises to 226% (KP14) and 20% (BGN) of the linear ceiling and the gap to +0.181 and +0.026, while GS21's mirrored ladder (`gvy25s`) gains room (2% to 11%) and no gap (−0.012). More priced factors (4) help only where each factor carries its own interaction: in the reduced form, factors that load on products of partly hidden states give room and a gap that rise with the number of factors and with the noise hiding the states, to +0.29 at ten factors, while factors that load linearly on observed characteristics give nothing at any K. The structural two-state economies (appendix) add no interaction: a second state through the same channel dilutes the exposure factor (KP14's room falls from 24% to 15%), and the designs that raise BGN's and GS21's gaps do so by giving each type its own state's premium, time variation in a sort rather than a second product.

Economically: every factor model is firm exposure times aggregate shock, so the interaction that matters is not with the aggregate state but within the exposure. A capturable gap needs a firm's exposure to a priced shock to be the product of two things that different characteristics reveal, "what the firm is exposed to" (its type, or its productivity) times "how levered the firm's claim is to it" (growth options in KP14, the project portfolio in BGN, the coupon in GS21). What no linear portfolio holds is the part of the exposure orthogonal to every linear function of the characteristics, and with one priced factor that is all or nothing: a linear portfolio with any covariance with the exposure mimics the factor and takes its whole Sharpe (the reduced form's sign ladder, Section 6). A type loading with a nonzero mean gives a sort that covariance, through the level compensation and the claim's leverage. That is why every positive-ladder structural gap is modest; why the GS21 economy in which leverage acts through the equity claim carries an interaction as large as KP14's (interaction SR 0.33) yet a near-zero gap, on a linear base of 0.81 (Section 6); and why the imposed GS21 product (productivity times coupon, mean zero), the reduced form's latent products, and KP14's mirrored type ladder (`vyg25s`: loadings ±0.02 to ±0.14, mean zero, hidden by the level compensation) reach gaps of +0.02 to +0.3, +0.18 in vyg25s: the linear span has nothing of theirs to hold.

The six economies and their stories. One per model and kind, each with the reading we attach to it; the constructions and numbers are in Sections 3b and 3c. Fair gaps are monthly Sharpe differences, ten seeds each.

slot	economy	fair gap	the story
KP14, non-monotone	`vyhx3`	+0.212 (t 12.5, ten of ten)	Five size classes of firms. The smallest and the largest load on the aggregate cash-flow shock with one sign and the mid-sized with the other, so exposure is hump-shaped in size and a sort on size cannot hold a premium that rises and falls along it. In KP14 the unexposed middle firms earn the most, so no characteristic ranks the exposure.
BGN, non-monotone	`bvyhx4` (and `bvyu5r`)	+0.119 (t 10.9, ten of ten); +0.0295 (t 3.1, nine of ten)	The same size hump in BGN. The one-signed U has the plainer story and reaches the bar at a loading of 0.35: the smallest firms are exposed through financing constraints, the largest through operating leverage, mid-sized firms through neither.
GS21, non-monotone	`gvyzzh10`	+0.080 (t 4.6, ten of ten)	A firm's exposure is quadratic in its own productivity. Firms of middling productivity bear the shock; the most productive have the margin to absorb it and the least productive have already shed the capacity that carries it. The premium is hump-shaped across ROE.
KP14, interaction	`vygelbcs`	+0.091 (t 12.3, ten of ten)	A priced shock with industry winners and losers, the sign carried by the characteristics only through size, times the firm's own profitability deviation: a dispersion shock, in which an exposed industry's firms above normal profitability pull further ahead and those below fall further behind. ROE carries the magnitude, size part of the sign, no single characteristic both.
BGN, interaction	`bvychs35`	+0.071 (t 5.8, ten of ten)	Industry winners and losers again, with the amplifier generated by the model: operating assets bear the shock and growth options largely do not, so firms with more of their value in assets in place move more.
GS21, interaction	`gvysz2`	+0.059 (t 11.5, ten of ten)	Industry winners and losers times the firm's productivity: a reallocation shock, in which the productive firms of an industry on the losing side gain share from the weak. Leverage, the more natural amplifier, fails here because fixed-coupon leverage types are a few observed groups the linear model prices directly.

The baselines every group is read against:

Table 1: Baselines. The three models as published, one type each: the floor every group is read against.

model	economy	$SR_{\max}$	market share	linear ceiling	room	room % of linear	interaction SR	DKKM	FMR	FF	linear	CAPM	fair gap	DKKM ⊥	+
BGN	`bgnbase` – baseline, no priced state beyond the rate	0.301	54%	0.250	+0.027	11	0.12	0.238	0.213	0.208	0.238	0.162	-0.0001 (t -0.0)	+0.00	6/10
GS21	`gsbase` – baseline	0.293	95%	0.291	+0.001	0	0.03	0.283	0.260	0.258	0.286	0.278	-0.0028 (t -4.3)	-0.03	2/10
KP14	`kpbase` – baseline, one type, no priced state	0.229	35%	0.221	+0.004	2	0.04	0.208	0.193	0.195	0.209	0.078	-0.0010 (t -0.6)	-0.00	6/10

Notation. BGN (Berk, Green and Naik 1999): a firm is a portfolio of live projects; project j has systematic risk β_j, drawn from an exponential distribution, and lognormal cash flows C_j; it survives each month with probability π; a new project arrives each month and is taken if its NPV is positive; the pricing kernel prices the market shock at σ_z and the Vasicek short rate r_t. GS21 (Gomes and Schmid 2021): output is $e^{x_t+z_{it}}k_{it}$ with aggregate productivity x_t and idiosyncratic productivity z_{it} both AR(1); firms choose the coupon b_{it}, invest lumpily at a random cost, and may default; the kernel prices x-shocks at a constant $\gamma_x = 0.5$. KP14 (Kogan and Papanikolaou 2014): firm value is assets in place plus growth options; project cash flows are $x_t u_{it} \varepsilon_{jt} K_j^{\alpha}$, new projects arrive at the firm-specific rate λ_{it}, and the kernel prices disembodied productivity x at γ_x and embodied technology z at γ_z.

The added priced state y. Every economy with a y adds one stationary state with unit stationary variance and mean reversion $\kappa_y = 0.35$ per year, $y_{t+1} = \rho y_t + \sqrt{1-\rho^2}\,\eta_{t+1}$ with $\rho = e^{-\kappa_y/12}$ (an Ornstein–Uhlenbeck process in KP14, a Tauchen chain in GS21 and in BGN's chain economies). It is priced: the kernel is multiplied by $e^{-\lambda_y \eta_{t+1} - \lambda_y^2/2}$ with $\lambda_y = \gamma_v \sqrt{1/12}$, and $\gamma_v = 2.5$ unless stated. "Types" are equal shares of firms differing only in their exposure b_f to y. "Level-compensated" means type f's cash flows are also scaled by a constant $\theta_f = (V_0/V_f)^{1.2}$, the ratio of a unit project's value for type 0 and type f at $y = 0$, so that value at $y = 0$ does not reveal the type (KP14's ω device).

The baselines.

- `bgnbase`: BGN as above, one type. Expected returns come from the value-weighted β_j of assets in place and from the growth option's exposure to r_t.
- `gsbase`: GS21 as above, Table I calibration, one type.
- `kpbase`: KP14 as above, one type, no y.

2 A different distribution of shocks: a crash-only factor

A cash-flow factor y (unit variance, half-life two years) multiplies every firm's cash flows as e^{by}, with one exposure b per model: 0.035 in BGN's project cash flows, $0.10\,\tilde{z}$ in GS21's output (the product economy's exposure without the coupon), 0.07 in KP14's project cash flows. Its Gaussian shocks carry no price. Its only priced event is a crash: each month, with probability 1%, y jumps to -3.5 standard deviations, and the kernel weights that transition by e^2 (about 7% a month risk-neutral); y then recovers at its ordinary mean reversion. Against the single-type economy of group 3a, the only change is that the Gaussian price is replaced by the crash.

Table 2: Group 1. A different distribution of shocks: a crash-only factor. Every firm's cash flows load on a factor y whose only priced shock is a rare crash (1% a month to -3.5 sd, kernel weight e^2).

model	economy	$\mathrm{SR}_{\max}$	market share	linear ceiling	room	room % of linear	interaction SR	DKKM	FMR	FF	linear	CAPM	fair gap	DKKM $\perp$	+
BGN	`bvy1j` – crash-only factor: $e^{0.035y}$, the same crash	0.324	50%	0.278	+0.033	12	0.14	0.265	0.239	0.235	0.262	0.160	+0.0003 (t 0.1)	+0.01	6/10
GS21	`gvyz10j` – crash-only factor on $0.10\,\tilde{z}\,y$: y unpriced except the 1%/month jump to -3.5 (weight e^2)	0.386	82%	0.374	+0.011	3	0.09	0.358	0.340	0.333	0.362	0.316	-0.0046 (t -3.5)	-0.05	1/10
KP14	`vy1j` – crash-only factor: $e^{0.07y}$, y unpriced except a 1%/month jump to -3.5 (weight e^2)	0.336	21%	0.277	+0.012	4	0.08	0.259	0.241	0.249	0.265	0.066	-0.0057 (t -2.3)	-0.04	2/10

Verdict. A crash changes the distribution of a shock every firm bears in the same proportion. The cross-section of premia it creates is a sort on the one firm state that scales the exposure, so the linear ceiling takes nearly all of it: room is 12% of the linear ceiling in BGN, 3% in GS21 and 4% in KP14, and the fair gap is zero or negative in all three (+0.000, −0.005, −0.006). Adding the same crash to the interaction economies instead (Appendix A1) leaves their room and gap at the control's in BGN and GS21 and halves KP14's.

What changes. The 3a economy with the Gaussian price of y set to zero ($\gamma_v = 0$) and a crash.

- `bvy1j`, `gvyz10j`, `vy1j`: `bvy1ch`, `gvyz10` and `vy1` with $\gamma_v = 0$ and, each month with probability 1% whatever y_t, a jump of y_{t+1} to -3.5 (the lowest node), weighted e^2 by the kernel (a risk-neutral probability of about 7% a month before renormalization). In KP14 the jump is a Poisson event at intensity 0.12 a year under P and $0.12\,e^2$ under Q.

3 A different price of risk: quadratic in the model's own state

The price of each model's own aggregate risk is multiplied by $1 + 0.25\,s^2$, capped at 3, where s is a standardized aggregate state: GS21's productivity x, and for KP14, whose own states are random walks, an unpriced stationary state y that no firm is exposed to. The multiplier never falls below 1 because KP14's growth-option discount turns negative once the price of x-risk drops below about 0.82 of its value. In BGN the price of the market shock σ_z is multiplied by the same function of the standardized rate, binned to seven equiprobable levels of the rate (1.00, 1.03, 1.17 and 1.68 from the middle outward) so that project values keep a closed form; the date-k cash flow is priced at the bin of the rate one period before it pays.

Table 3: Group 2. A quadratic price of the model's own risk. The price of the baseline's aggregate risk becomes $1 + 0.25s^2$ times its value, capped at 3, s a standardized aggregate state.

model	economy	$\mathrm{SR}_{\max}$	market share	linear ceiling	room	room % of linear	interaction SR	DKKM	FMR	FF	linear	CAPM	fair gap	DKKM $\perp$	+
BGN	`bq25` – $\sigma_z \times \min(1 + 0.25\, s^2, 3)$, s the standardized rate, 7 bins	0.411	36%	0.346	+0.049	15	0.18	0.339	0.303	0.293	0.335	0.149	+0.0034 (t 0.7)	+0.00	4/10
GS21	`gq25` – $\gamma_x \times \min(1 + 0.25\, \tilde{x}^2, 3)$	0.345	95%	0.343	+0.002	0	0.03	0.334	0.316	0.309	0.337	0.326	-0.0036 (t -5.5)	-0.05	0/10
KP14	`kpq` – $\gamma_x \times \min(1 + 0.25\, y^2, 3)$, y unpriced	0.260	53%	0.251	+0.003	1	0.04	0.240	0.226	0.225	0.242	0.137	-0.0021 (t -1.5)	-0.02	4/10

Verdict. A price of risk that moves with the aggregate state rotates the premium across states and leaves it linear in the ranks within one. The quadratic price gives room under 1% of the linear ceiling in GS21 and KP14, and 14% in BGN, where the rate is itself a priced state so a price that rises with its distance from the mean rotates the premium with it; the fair gap is zero in all three (+0.003, −0.004, −0.002). The regime prices tried earlier (appendix) do the same, and BGN's fast regime gave the group's one small positive number (+0.008), which a ranks-times-state linear benchmark takes.

What changes.

- `bq25`: `bgnbase` with $\sigma_z \to \sigma_z\, g(r_t)$, $g(r) = \min(1 + 0.25\,((r - \bar{r})/\mathrm{sd}_r)^2, 3)$ binned to seven equiprobable bins of the stationary rate (g = 1.00, 1.03, 1.17, 1.68 from the middle outward); the date-k cash flow is priced at the bin of r_{t+k-1}, through seven bases $D_j(r)$ in place of the regime's two.
- `gq25`: `gsbase` with $\gamma_x \to \gamma_x (1 + 0.25\, \tilde{x}^2)$ capped at 1.5, $\tilde{x} = x/\mathrm{sd}(x)$.
- `kpq`: `kpbase` with $\gamma_x \to \gamma_x \min(1 + 0.25\, y^2, 3)$, y the unpriced stationary state of `kpgam` ($\kappa = 0.35$, unit variance); no firm is exposed to y.

4 Different exposures to aggregate risk

4.1 Exposure that a single characteristic captures

The priced state y is added with one type: every firm's cash flows load on it with the same exponent (the group 1 exposures, now with y's Gaussian shocks priced at 2.5). A firm's exposure to y then varies across firms only through one firm state: the mix of assets in place and growth option in BGN, productivity in GS21 (the product economy's exposure without the coupon), the growth-option share in KP14.

Table 4: Group 3a. Exposure captured by a single characteristic. A priced state y with one type: a firm's exposure varies only through one firm state (project mix, operating leverage, growth-option share).

model	economy	$\mathrm{SR}_{\max}$	market share	linear ceiling	room	room % of linear	interaction SR	DKKM	FMR	FF	linear	CAPM	fair gap	DKKM $\perp$	+
BGN	`bvy1ch` – one type, project cash flows $\times e^{0.035y}$	0.336	54%	0.297	+0.029	10	0.13	0.283	0.252	0.255	0.281	0.181	+0.0017 (t 0.8)	+0.02	7/10
GS21	`gvyz10` – one type, output $\times e^{0.10\,\tilde{z}\,y}$: the product economy without the coupon	0.854	50%	0.800	+0.043	5	0.27	0.793	0.805	0.742	0.808	0.421	-0.0164 (t -6.5)	-0.15	1/10
KP14	`vy1` – one type, cash flows $\times e^{0.07y}$	0.466	12%	0.369	+0.019	5	0.12	0.355	0.335	0.340	0.360	0.047	-0.0050 (t -2.9)	-0.05	1/10

Verdict. A monotone function of one firm state, however nonlinear, is linear in the rank of the characteristic that reveals it. Room is 10% of the linear ceiling in BGN and 5% in GS21 and KP14, and the fair gap is zero or negative in all three (+0.002, −0.017, −0.005); in GS21 the productivity exposure alone carries a large premium ($\mathrm{SR}_{\max}$ 0.85) that the linear ceiling captures almost entirely. The same holds for every other single-characteristic wiring tried (Appendix A3), including ones with far more room: room is necessary, not sufficient. The absence of room is not an artefact of the rank standardization DKKM applies to its random features after the sines and cosines: recomputing the ceilings on the saved panels with the raw features (two seeds per economy, `variants/mlp/room_raw.py`) gives room of 4% / 6% / 8% of the linear ceiling in `vy1` / `gvyz10` / `bvy1ch` against 5% / 6% / 8% with the standardization, and 19% / 62% / 10% against 26% / 61% / 11% in the 3b economies; the quadratic and bin ceilings, which never standardize, say the same.

What changes.

- `bvy1ch`: `bgnbase` plus y (21-node chain), one type; project cash flows $C_j e^{0.035\, y_t}$.
- `gvyz10`: `gsbase` plus y (9-node chain over ±3.5 sd), one type; output $e^{x_t + z_{it} + 0.10\, \tilde{z}_{it}\, y_t} k_{it}$, $\tilde{z} = z/\mathrm{sd}(z)$: the product economy's exposure without the coupon.
- `vy1`: `kpbase` plus y, one type; project cash flows $e^{0.07\, y_t} x_t u_{it} \varepsilon_{jt} K_j^{\alpha}$.

4.2 Exposure as an interaction of several characteristics

The same y with the exposure a product of two firm-level factors, one carrying a sign that no monotone sort on a characteristic reads. The construction every model carries: a mirrored set of types, the winners and losers of the priced shock (an industry attribute the characteristics do not include and reveal only in part, through size), each level-compensated at $y = 0$ so that value does not reveal it, times a firm-level amplifier. The amplifier differs: in BGN the share of value in assets in place (model-generated), in KP14 the profitability deviation (in the cash flow), in GS21 the productivity (in the production function). The original three-type economies, the levers tried, and the leverage version of GS21's amplifier are in Section 6.

Table 5: Group 3b. Exposure as an interaction of characteristics. One construction in every model: a mirrored set of types, the winners and losers of the priced shock, level-compensated at $y = 0$ and revealed only in part through size, times a firm-level amplifier – the assets-in-place share in BGN (bvychs35), the profitability deviation in KP14 (vygelbcs, bounded exposure, price 4.0, σ_ϵ 0.3), the productivity in GS21 (gvysz2).

model	economy	$\mathrm{SR}_{\max}$	market share	linear ceiling	room	room % of linear	interaction SR	DKKM	FMR	FF	linear	CAPM	fair gap	DKKM ⊥	+
BGN	`bvychs35` – bvychs at bv_comp 3.5	0.508	31%	0.294	+0.189	67	0.38	0.403	0.315	0.309	0.303	0.156	+0.0712 (t 5.8)	+0.22	10/10
GS21	`gvysz2` – type sign × productivity $\tilde{z}$: mirrored types ±0.15 / 0.3 / 0.5, gs_level ladder 0.5..2.0, no compensation	0.847	17%	0.603	+0.208	37	0.54	0.678	0.536	0.488	0.619	0.151	+0.0590 (t 11.5)	+0.27	10/10
KP14	`vygelbcs` – bounded exposure (cap 2.5), mirrored types ±0.3 / 0.7 / 1.2, level ladder 0.5..2.0, γ_v 4.0, σ_ϵ 0.3	0.858	7%	0.297	+0.256	90	0.46	0.388	0.251	0.288	0.282	0.052	+0.0913 (t 12.3)	+0.25	10/10

Why a product has room. A linear model in the ranks fits additive functions of the characteristics; the product of two factors each mean zero across firms has no additive part, so no linear portfolio is long the exposure, while random features span products. The sign must be hidden: one-sided versions are a sort on the amplifier's characteristic with a type-dependent slope, room 10–27%, no gap. The $y = 0$ compensation keeps value from revealing the sign; how visible the type is through size does not matter for room.

BGN: `bvychs35`, +0.071. Six mirrored types, loadings ±0.01, ±0.035, ±0.07 on project cash flows, compensation exponent 3.5 (the type becomes a size characteristic, room 67%). The amplifier is claim leverage: operating assets bear the shock, growth options largely do not. Fair gap +0.071 (t 5.8, ten of ten; DKKM 0.403 against 0.303); the compensation ladder 2.5 / 3.0 / 3.5 / 4.0 gives +0.048 / +0.063 / +0.071 / +0.059.

KP14: `vygelbcs`, +0.091. The exposure on the profitability deviation, cash flow $xK^{\alpha}\theta u\,[1+(\varepsilon-1)h_b(y)]$, six mirrored types ±0.3, ±0.7, ±1.2 revealed by a level ladder 0.5–2.0, price 4.0, σ_ε 0.3; $h_b(y) = c/(1+(c-1)e^{-cby/(c-1)})$ with $c = 2.5$ replaces e^{by}, which turns project values negative for $b \leq -0.2$. Room 58–119% (0.15–0.30 in Sharpe), fair gap +0.091 (t 12.3, ten of ten; +0.041 to +0.131; DKKM 0.388 against 0.297): the type's sign read from size, the firm's deviation from ROE, their product spanned by no linear function of the characteristics.

GS21: `gvysz2`, +0.059. The exposure gs_bv$_f\,\tilde{z}$, six mirrored types ±0.15, ±0.3, ±0.5 on the new gs_level ladder 0.5–2.0, no compensation (the exposure is mean zero in $\tilde{z}$). Room 22–66%, fair gap +0.059 (t 11.5, ten of ten; DKKM 0.678 against 0.619); the smaller ladder (`gvysz1`) +0.053 (t 6.1). The type premium across the ladder, −0.30, +0.01, +0.11, +0.51, +0.42, +0.24% a month, is non-monotone in b. The imposed product without types, productivity × coupon (`gvyp15`, +0.078), is the same mechanism without the industry sign; leverage as the amplifier fails as a type (Section 6).

What changes.

- `bvychs35`: `bvych0` with the ladder mirrored (six types, $b_f = \pm0.01, \pm0.035, \pm0.07$) and the compensation exponent 3.5; chain y, price 2.5.
- `vygelbcs`: `vy1`'s y on the profitability deviation, cash flow $xK^{\alpha}\theta u\,[1+(\varepsilon-1)h_b(y)]$, bounded h_b (cap 2.5); six mirrored types $\pm0.3, \pm0.7, \pm1.2$ on the level ladder 0.5–2.0; price 4.0; σ_ε 0.3.
- `gvysz2`: `gvy25s`'s chassis with the exposure gs_bv$_f\,\tilde{z}$, six mirrored types $\pm0.15, \pm0.3, \pm0.5$ on the gs_level ladder 0.5–2.0, bv_comp 0.

4.3 Exposure that is non-monotone in a characteristic

A monotone exposure in one firm state is linear in that state's rank; an exposure that rises and then falls in a characteristic is not, and a linear model fits a line through the hump or the U. The construction, the same in every model that allows it: five equal types on a mild visible level ladder (the cash-flow or production level × 0.8, 0.9, 1.0, 1.1, 1.2; BGN's and KP14's type_level, GS21's new gs_level), each level-compensated at $y = 0$ (bv_comp 1.0), the loading hump-shaped, $b = -c, 0, +c, 0, -c$, or U-shaped, $+c$ at the ends, price 2.5. The hump's story is the thinner of the two shapes, the tails of the size distribution on one side of the shock and the middle on the other; the U, with the plainer story (small firms exposed through financing constraints, large firms through operating leverage), reaches the bar only in BGN and only just, and is in Section 6 with the leverage U.

Table 6: Group 3c. Exposure non-monotone in a characteristic. A hump-shaped loading ($b = -c, 0, +c, 0, -c$) across five equal types on a mild visible level ladder (0.8 to 1.2), level-compensated at $y = 0$, in BGN (bvyhx4, $c = 0.28$) and KP14 (vyhx3, $c = 0.21$); GS21, where the ladder is priced linearly, carries the exposure quadratic in the firm's own productivity (gvyzzh10).

model	economy	SR$_{\max}$	market share	linear ceiling	room	room % of linear	interaction SR	DKKM	FMR	FF	linear	CAPM	fair gap	DKKM ⊥	+
BGN	`bvyhx4` – hump-shaped loading across a mild level ladder 0.8..1.2: b = -0.28, +0.00, +0.28, +0.00, -0.28	0.642	25%	0.357	+0.260	73	0.50	0.476	0.341	0.293	0.344	0.159	+0.1186 (t 10.9)	+0.31	10/10
KP14	`vyhx3` – hump-shaped loading across a mild level ladder 0.8..1.2: b = -0.21, +0.00, +0.21, +0.00, -0.21	0.632	7%	0.207	+0.377	197	0.54	0.431	0.217	0.136	0.178	0.031	+0.2123 (t 12.5)	+0.37	10/10
GS21	`gvyzzh10` – hump-shaped exposure in the firm's own productivity: $-0.1\,(\tilde{z}^2-1)$, one type	0.639	25%	0.364	+0.250	70	0.49	0.480	0.325	0.255	0.400	0.152	+0.0795 (t 4.6)	+0.25	10/10

The hump across size: KP14 and BGN carry it, GS21 does not. BGN's ladder rises with the loading: room 9–13% at ±0.14, 37–40% and +0.035 (t 3.9, ten of ten) at ±0.21, 59–77% and +0.119 (t 10.9, ten of ten) at ±0.28 (`bvyhx2`, `bvyhx3`, `bvyhx4`). KP14's is the largest gap of this kind in the campaign: at ±0.21 (`vyhx3`) room 154–197% of a linear ceiling of 0.21–0.23 and a fair gap of +0.212 (t 12.5, ten of ten; +0.134 to +0.306; DKKM 0.431 against 0.218), at ±0.28 (`vyhx4`) room 123% and +0.166 (t 10.6, ten of ten). The type premium across the ladder is 0.36, 0.86, 0.68, 0.88, 0.35% a month (vyhx3, seed 0): the unexposed middle types earn the most, the market beta rises from 0.73 to 1.36

with b while the y channel pulls the other way, so the premium is non-monotone in the exposure itself and no characteristic that reveals the type can rank it. GS21's ladder gives nothing (`gvyhx3`: room 9–31%, +0.002, t 0.9): its premium is monotone in the exposure (−0.50, +0.26, +1.47, +0.25, −0.56% a month), and a production level moves capital, value and profitability together, so book-to-market and ROE rank the exposure and the linear ceiling (0.4–0.9) holds it.

The quadratic in productivity. GS21's working non-monotone exposure: one type whose exposure to y is quadratic in its own standardized productivity, $c\,(\tilde{z}^2 - 1)$, mean zero across firms (`gvyzzh10`, `gvyzz20`: $c = -0.1$, $+0.2$). The premium across ROE quintiles is a hump (−0.24, +0.58, +0.41, +0.34, −0.54% a month) or a U, room 55–81% and 25–65%, fair gap +0.080 (t 4.6, ten of ten) and +0.059 (t 6.3, ten of ten); the smaller U ($c = +0.1$) has the room and not the gap, the leverage characteristic revealing it. The story: firms of middling productivity are the most exposed, the most productive having the margin to absorb the shock, the least productive having already shed the capacity that bears it.

The three economies and their stories. KP14 `vyhx3` (+0.212) and BGN `bvyhx4` (+0.119), the hump across size, with BGN's U in size at the bar at 0.35 (`bvyu5r`, +0.0295); GS21 `gvyzzh10` (+0.080), the quadratic in productivity. In each, the exposure is a function of the characteristics that no linear combination of them spans, and the premium is non-monotone in whatever the characteristics reveal about the type; where the premium is monotone in a revealed quantity (the fixed-coupon leverage types, the GS21 size ladder, the one-signed U) the linear model holds it however curved the exposure is.

What changes.

- `bvyhx4`: `bvy1ch` with five types on the level ladder θ_f 0.8–1.2, loadings $-0.28, 0, +0.28, 0, -0.28$, bv_comp 1.0, chain, price 2.5.
- `vyhx3`: the same ladder in KP14 (type_level), loadings $-0.21, 0, +0.21, 0, -0.21$, bv_comp 1.0, price 2.5.
- `gvyzzh10`: `gvy25s`'s chassis with one type whose exposure is $-0.1\,(\tilde{z}^2 - 1)$ (gs_bzz 1, gs_bzl 0), no compensation, price 2.5.

5 Larger factor models

More priced factors are tested where the structure is controlled: the reduced form below adds K priced factors to a market factor, with the total price of risk held across K, and varies only whether the factors load linearly on the observed characteristics or on products of partly hidden states (setup below). The structural version, a second priced state in each of the three models, is in the appendix (Table 9).

Factors that load linearly on observed characteristics leave no room at any K and a negative gap; factors that load on products of partly hidden states give room and a gap that grow monotonically with K and with the observation noise that hides the states. Curvature is the third kind of loading in the table: squares of observed characteristics (`kfh_sq`, `kfo_sq`) and the square of one partly hidden latent (`lsqb`). Bounded loadings and the sign ladder are in Section 6.

Bounded loadings behave like linear ones: no room at $K = 1$ and 4, a little at $K = 16$ (12%), and a negative gap throughout, because tanh of a single index is monotone in that index and nearly linear in its rank. Pure curvature of an observed characteristic gives room of about five times the linear ceiling and a gap of about +0.5 at $K = 1$. In the unconstrained square rows the gap then falls to +0.12 and +0.14 at $K = 4$ and 16 while the room does not move: the squared index contains size^2, Fama–French's legs are value-weighted and market value is exp(size), so those legs hold a size^2 exposure mechanically

Table 7: Reduced form: priced factors loading linearly on observed characteristics, on products of latent states, and on curvature (squares of observed characteristics, or of one partly hidden latent). Five seeds per cell, the total price of risk fixed across K (per-factor price $\propto 1/\sqrt{K}$); K = number of priced factors; the lfh rows load products of pairs of the five latent processes, observed with noise of sd 0 (n0) or 1 (n1).

configuration	$\mathrm{SR}_{\max}$	market share	linear ceiling	room	room % of linear	interaction SR	DKKM	FMR	FF	linear	CAPM	fair gap	+
kfh_lin_1 – K=1 factor on observed characteristics, linear loadings	0.598	8	0.598	+0.000	0	0.01	0.590	0.584	0.581	0.595	0.041	-0.0067 (t -1.7)	0/5
kfh_lin_4 – K=4	0.592	8	0.588	+0.003	0	0.06	0.555	0.578	0.567	0.575	0.041	-0.0237 (t -3.1)	0/5
kfh_lin_16 – K=16	0.385	13	0.383	+0.001	0	0.03	0.352	0.357	0.353	0.365	0.041	-0.0128 (t -1.2)	2/5
lfh_L5_K1_n0 – K=1 product of latent states, no observation noise	0.598	8	0.208	+0.389	187	0.56	0.570	0.156	0.539	0.175	0.041	+0.0311 (t 4.0)	5/5
lfh_L5_K2_n0 – K=2	0.598	8	0.144	+0.453	314	0.58	0.564	0.111	0.520	0.115	0.041	+0.0443 (t 6.7)	5/5
lfh_L5_K3_n0 – K=3	0.598	8	0.106	+0.490	464	0.59	0.566	0.059	0.507	0.084	0.041	+0.0589 (t 9.0)	5/5
lfh_L5_K5_n0 – K=5	0.598	8	0.092	+0.502	549	0.59	0.543	0.048	0.423	0.064	0.041	+0.1193 (t 21.6)	5/5
lfh_L5_K10_n0 – K=10	0.598	8	0.089	+0.503	568	0.58	0.507	0.047	0.250	0.067	0.041	+0.2572 (t 39.9)	5/5
lfh_L5_K1_n1 – K=1 product of latent states, observation noise 1.0	0.598	8	0.167	+0.416	249	0.56	0.526	0.096	0.433	0.137	0.041	+0.0930 (t 3.8)	5/5
lfh_L5_K2_n1 – K=2	0.598	8	0.111	+0.469	423	0.57	0.523	0.072	0.420	0.082	0.041	+0.1025 (t 6.8)	5/5
lfh_L5_K3_n1 – K=3	0.598	8	0.091	+0.486	537	0.57	0.528	0.057	0.391	0.073	0.041	+0.1371 (t 10.4)	5/5
lfh_L5_K5_n1 – K=5	0.598	8	0.075	+0.496	664	0.57	0.496	0.038	0.309	0.060	0.041	+0.1866 (t 26.3)	5/5
lfh_L5_K10_n1 – K=10	0.598	8	0.079	+0.479	607	0.55	0.456	0.052	0.171	0.077	0.041	+0.2854 (t 24.8)	5/5
kfh_sq_1 – K=1, curvature: squares of observed characteristics	0.600	8	0.080	+0.519	647	0.59	0.582	0.032	0.046	0.059	0.041	+0.5161 (t 57.3)	5/5
kfh_sq_4 – K=4	0.601	8	0.097	+0.501	515	0.59	0.573	0.042	0.458	0.074	0.041	+0.1155 (t 22.6)	5/5
kfh_sq_16 – K=16	0.601	8	0.076	+0.523	687	0.59	0.571	0.036	0.429	0.052	0.041	+0.1429 (t 14.3)	5/5
kfo_sq_1 – K=1, squares, directions orthogonal to size	0.600	8	0.080	+0.519	648	0.59	0.582	0.032	0.048	0.059	0.041	+0.5163 (t 57.7)	5/5
kfo_sq_4 – K=4	0.601	8	0.092	+0.507	552	0.59	0.575	0.058	0.061	0.081	0.041	+0.4833 (t 25.4)	5/5
kfo_sq_16 – K=16	0.602	8	0.079	+0.520	656	0.59	0.575	0.015	0.031	0.053	0.041	+0.5030 (t 21.3)	5/5
lsqb_a0_n0 – K=1 square of the book-to-market latent, mean zero, no noise	0.598	8	0.199	+0.398	200	0.56	0.564	0.142	0.125	0.173	0.041	+0.3662 (t 11.0)	5/5
lsqb_a0_n1 – noise 1.0	0.598	8	0.190	+0.392	206	0.55	0.515	0.120	0.092	0.146	0.041	+0.3403 (t 6.5)	5/5

and Fama–French's Sharpe jumps from 0.05 to 0.46. The `kfo` rows remove the size coordinate from every direction and the gap returns to about +0.5 at every K. Curvature of an observed characteristic is therefore captured by random features and by nothing linear, and the square ladder's apparent fall with K was an artefact of value weighting. A factor on the square of one partly hidden latent is the same result one level down: the square of the book-to-market latent, mean zero (`lsqb_a0_n0`), has room +0.40 and a fair gap of +0.37 (t 11; +0.34 with observation noise 1.0), while the same square on the size latent has the room and no fair gap, because market value is e^{size} and the value-weighted market inside the Fama–French set holds size^2; any shift of the latent's mean adds a linear part and sends room to 0.002 (Section 6).

The reduced form. A simulated panel (`variants/reduced/rf_model.py`) of $N = 500$ firms over $T = 500$ months after a 100-month burn-in, five seeds per cell, run through the same oracle ceilings and the same estimators as the structural economies at the 360-month window; its true conditional moments are known in closed form. Notation: i indexes firms, t months, $j = 1, \ldots, 5$ the characteristics (size, book-to-market, asset growth, profitability, momentum), $k = 1, \ldots, K$ the extra priced factors.

- **Latent processes.** Each firm carries five independent AR(1) processes $z_{ijt} = \rho_j z_{ij,t-1} + \sqrt{1-\rho_j^2}\,\eta_{ijt}$, η standard normal, $\rho = (0.99, 0.98, 0.90, 0.95, 0.70)$, so each has unit stationary variance.

- **Observed characteristics.** $X_{ijt} = z_{ijt}$ in the `kfh` and `kfo` ladders. In the `lfh` ladder the five processes are the firm's true states and each characteristic is its state plus observation noise, $X_{ijt} = z_{ijt} + \sigma_\nu \nu_{ijt}$, with ν_{ijt} an independent AR(1) with the same ρ_j and unit variance. The ladder's one noise parameter is σ_ν: 0 (the states are observed exactly, rows `n0`) or 1 (noise as large as the state itself, rows `n1`). $\tilde{X}_{it}$ is X_{it} standardized across firms at t. The oracle's bases and the estimators see the ranks of X_{it}, as everywhere in the report.

- **Returns.** $r_{i,t+1} = B_{it}\lambda + B_{it} f_{t+1} + \varepsilon_{i,t+1}$, with $f_{t+1} \sim N(0, \sigma_f^2 I_{K+1})$, $\sigma_f = 0.04$, and $\varepsilon_{i,t+1} \sim N(0, \sigma_\varepsilon^2)$, $\sigma_\varepsilon = 0.10$, all independent. B_{it} is the $1 \times (K+1)$ row of firm i's loadings and λ the $(K+1)$-vector of constant prices of risk. The true premium is $\mu_{it} = B_{it}\lambda$ and the true covariance $\Sigma_t = \sigma_f^2 B_t B_t' + \sigma_\varepsilon^2 I$.

- **The market factor (column 0).** $B_{i0t} = 1 + a'\tilde{X}_{it}$ with $a = (0, 0.3, 0, 0.3, 0.3)$ on (size, bm, agr, roe, mom), and $\lambda_0 = 0.002$, a Sharpe of $\lambda_0/\sigma_f = 0.05$. This loading is linear in the characteristics, so every method holds it.

- **The K extra factors (columns $k = 1, \ldots, K$).** $\lambda_k = 0.024/\sqrt{K}$ for every k, so each has Sharpe $0.6/\sqrt{K}$ and the total price of risk is the same at every K. Every extra loading is centered across firms each month ($\bar{\beta}_{kt}$ is the cross-sectional mean), so the equal-weighted market carries none of the extra premia. The ladders differ only in the loading:
  - ***`kfh`*** (loadings on observed characteristics): $B_{ikt} = h(w_k' \tilde{X}_{it}) - \bar{\beta}_{kt}$, with w_k a fixed unit vector in $\mathbb{R}^5$ drawn once and reused in every seed, and $h(u) = u$ in the main table (lin); the appendix adds $h(u) = \tanh(u)$ (tanh) and $h(u) = u^2$ (sq). Exposure is a function of what the characteristics show directly: GS21's structure.
  - ***`kfo`*** (appendix): `kfh` with the size coordinate of each w_k set to zero before normalizing, so no extra factor loads on size.
  - ***`lfh`*** (loadings on products of latent states): $B_{ikt} = z_{ia_kt}\, z_{ib_kt} - \bar{\beta}_{kt}$, where (a_k, b_k) is the k-th pair in a fixed enumeration of the ten pairs of the five processes (so $K \leq 10$). The loading is a product of two true states that the characteristics reveal only up to the observation noise σ_ν: KP14's and BGN's structure, a type times a within-firm state.

The ladders therefore vary one thing at a time: K; whether the loading is a function of the observed characteristics or a product of partly hidden states; and, within `lfh`, how much noise hides the states. In the `kfh` linear and tanh ladders $\text{SR}_{\max}$ falls at $K = 16$ because sixteen directions over five characteristics are collinear and their loadings partly cancel; in the `lfh` ladder $\text{SR}_{\max}$ stays near 0.6 at every K.

6 Additional economies

Everything else that was run. The table has the main tables' columns; the bullets describe each economy against its parent (notation as in Section 1). Three groups of them answer side questions: the crash added to the interaction economies instead of to a single-type economy (`bvyj25`, `gvyp10j`, `vyj25`), the GS21 product ladder (`gvyp05` to `gvyp20`, `gvyp10x`), and the earlier single-characteristic and price-of-risk wirings that the harmonized cells replace.

BGN

- `g0235f`: `bgnbase` with $\sigma_z \to g_s \sigma_z$, $s_t \in \{\text{calm}, \text{stress}\}$ a Markov chain, $g = 0.2$ calm and 3.5 stress; monthly switch probabilities 1/12 (calm to stress) and 1/6 (stress to calm), so stress is one third of months with a mean spell of 6 months.

- `g0235df`: the same g with switch probabilities 1/6 and 1/12: stress two thirds of months, mean spell 12 months.

- `bvych0`: `bvy1ch` with three types, $b_f \in \{0.01, 0.035, 0.07\}$, level-compensated ($\theta_f = (V_0/V_f)^{1.2}$). Exposure is b_f times the firm's mix of assets in place and growth option.

- `bvy25`: `bgnbase` plus y; type f's project cash flows are $\theta_f C_j e^{b_f y_t}$, $b_f \in \{0.01, 0.035, 0.07\}$, level-compensated. Exposure is b_f times the firm's mix of assets in place and growth option.

- `bvych41`: `bvych0` with y on a 41-node chain instead of 21.

- `bvy9t2`: `bvych0` with a second independent 21-node chain and nine types, a pair (f_1, f_2) drawn independently from $\{0.01, 0.035, 0.07\}$ for each chain, cash flows $C_j\, e^{b_{f_1} y_t + b_{f_2} y_t^{(2)}}$; price $2.5/\sqrt{2}$ on each: the nested nine-type design of `vy9t2` in BGN.

Table 8: Additional economies. Everything else that was run, by model, each described in the appendix against its parent; the RF rows are the reduced-form ladders with bounded (tanh), curvature (square) and size-orthogonal curvature (kfo) loadings, the sign ladder (lsh: the K=1 latent product with the factors' means shifted; lrt: a one-signed ratio), and the curvature-in-a-latent ladder (lsq / lsqb: the square of one partly hidden latent, size or book-to-market, with its mean shifted).

model	economy	$SR_{\max}$	market share	linear ceiling	room	room % of linear	interaction SR	DKKM	FMR	FF	linear	CAPM	fair gap	DKKM ⊥	+
BGN	`g0235f` – regime price ×0.2 / ×3.5, spells 4× faster	0.237	62%	0.163	+0.045	28	0.13	0.169	0.141	0.131	0.161	0.146	+0.0079 (t 4.8)	+0.05	10/10
BGN	`g0235df` – stress 2/3 of the time, 4× faster	0.489	17%	0.432	+0.036	9	0.18	0.423	0.392	0.351	0.428	0.077	-0.0051 (t -2.2)	-0.05	2/10
BGN	`bvy25` – three types on y, level-compensated (1.2)	0.601	41%	0.398	+0.062	16	0.23	0.399	0.344	0.367	0.384	0.248	+0.0148 (t 3.0)	+0.09	9/10
BGN	`bvych41` – bvych0 on a 41-node chain	0.565	38%	0.336	+0.052	15	0.19	0.328	0.279	0.291	0.320	0.215	+0.0086 (t 2.8)	+0.06	8/10
BGN	`bvyc0` – the same without compensation	0.475	41%	0.395	+0.066	18	0.22	0.386	0.321	0.347	0.376	0.195	+0.0104 (t 3.1)	+0.07	9/10
BGN	`bvyj25` – the chain with a priced crash (1%/month to $-y_{\max}$, priced e^2)	0.536	38%	0.315	+0.054	17	0.19	0.307	0.261	0.269	0.301	0.202	+0.0066 (t 2.8)	+0.05	9/10
BGN	`bvysg` – signed exposure ladder (−0.07 / 0 / +0.07)	0.504	28%	0.361	+0.120	35	0.31	0.318	0.305	0.273	0.324	0.141	-0.0199 (t -2.0)	-0.08	2/10
GS21	`g28` – countercyclical $\gamma(x)$, clipped	0.314	95%	0.312	+0.001	0	0.02	0.303	0.254	0.256	0.305	0.298	-0.0024 (t -4.3)	-0.04	0/10
GS21	`bx7` – two-state regime on γ, five types on x	0.327	94%	0.317	+0.007	2	0.07	0.309	0.282	0.271	0.311	0.306	-0.0017 (t -2.8)	-0.03	0/10
GS21	`gx7` – five types on productivity x	0.485	98%	0.474	+0.008	2	0.08	0.473	0.393	0.383	0.473	0.473	+0.0003 (t 1.6)	+0.01	5/10
GS21	`gvy1` – one type, output $\times e^{0.05y}$	0.636	95%	0.633	+0.002	0	0.05	0.626	0.608	0.598	0.629	0.602	-0.0030 (t -4.0)	-0.05	1/10
GS21	`gvy1j` – crash-only factor: $e^{0.05y}$, the same crash	0.379	96%	0.377	+0.001	0	0.03	0.372	0.350	0.341	0.374	0.363	-0.0022 (t -4.4)	-0.04	1/10
GS21	`gvy25` – production-level exposure to y, three types	0.725	86%	0.709	+0.014	2	0.14	0.688	0.682	0.666	0.695	0.625	-0.0070 (t -4.0)	-0.09	1/10
GS21	`gvys25` – signed fixed-cost exposure to y	0.619	69%	0.602	+0.014	2	0.13	0.496	0.490	0.426	0.499	0.420	-0.0032 (t -1.5)	-0.03	2/6
GS21	`gvyp05` – product $\tilde{z}\tilde{b}$ at 0.05	0.425	60%	0.284	+0.128	45	0.30	0.297	0.254	0.202	0.276	0.254	+0.0203 (t 7.4)	+0.11	10/10
GS21	`gvyp075` – product at 0.075	0.561	28%	0.261	+0.287	111	0.48	0.433	0.212	0.187	0.352	0.149	+0.0802 (t 9.4)	+0.25	10/10
GS21	`gvyp15` – product at 0.15	0.744	-25%	0.616	+0.119	22	0.39	0.665	0.349	0.516	0.612	0.192	+0.0528 (t 9.3)	+0.26	10/10
GS21	`gvyp20` – product at 0.20	0.778	-51%	0.729	+0.039	6	0.23	0.710	0.452	0.643	0.715	0.407	-0.0049 (t -1.1)	-0.04	4/10
GS21	`gvyp10x` – product at 0.10 plus a curvature loading on x	0.589	17%	0.365	+0.215	63	0.45	0.491	0.196	0.262	0.448	0.054	+0.0431 (t 7.6)	+0.20	10/10
GS21	`gvyp10j` – the product exposure with a priced crash	0.537	21%	0.328	+0.197	62	0.41	0.442	0.242	0.246	0.415	0.072	+0.0267 (t 3.8)	+0.14	10/10
GS21	`gvypk2` – two priced chains, two product types	0.478	27%	0.226	+0.212	103	0.37	0.309	0.168	0.153	0.254	0.105	+0.0559 (t 3.2)	+0.15	9/10
GS21	`gvyzl` – productivity exposure, three fixed-coupon leverage types (diluting)	0.865	50%	0.793	+0.059	8	0.31	0.804	0.793	0.743	0.799	0.426	+0.0049 (t 2.5)	+0.05	7/10
GS21	`gvyzlf` – productivity exposure, three leverage types, capital structure fixed for life	0.880	50%	0.808	+0.066	8	0.33	0.820	0.807	0.758	0.815	0.432	+0.0046 (t 2.0)	+0.05	7/10
GS21	`gvyzlw` – gvyzlf with coupons 0.05 / 0.45 / 0.85	0.782	52%	0.729	+0.048	7	0.27	0.722	0.730	0.676	0.738	0.404	-0.0165 (t -5.5)	-0.15	0/10
GS21	`gvyzl5` – gvyzlf with five leverage types	0.880	50%	0.812	+0.060	7	0.32	0.818	0.812	0.765	0.818	0.440	-0.0016 (t -0.7)	-0.03	4/10
GS21	`gvyzlb` – gvyzlf with the exposure doubled	1.066	45%	1.003	+0.056	7	0.31	0.996	0.968	0.944	0.985	0.490	+0.0114 (t 1.5)	+0.06	6/10
GS21	`gvyzlr` – gvyzlf with riskless debt	0.880	49%	0.808	+0.066	8	0.33	0.820	0.807	0.758	0.815	0.432	+0.0045 (t 2.0)	+0.05	7/10
GS21	`gvyzlp` – gvyzlf with a persistent y (κ_y 0.10)	0.475	76%	0.457	+0.017	4	0.13	0.448	0.428	0.418	0.451	0.355	-0.0035 (t -2.7)	-0.04	2/10
GS21	`gvyzlrp` – gvyzlf with riskless debt and a persistent y	0.475	76%	0.457	+0.017	4	0.13	0.448	0.428	0.419	0.451	0.356	-0.0035 (t -2.7)	-0.04	2/10
GS21	`gsbg1` – benchmark at γ_x 0.1	0.070	96%	0.070	+0.000	0	0.00	0.061	0.029	0.030	0.061	0.054	+0.0004 (t 0.9)	+0.00	6/10
GS21	`gvyz10g` – gvyz10 at γ_x 0.1	0.770	25%	0.707	+0.051	7	0.27	0.703	0.722	0.647	0.724	0.195	-0.0229 (t -6.5)	-0.18	0/10
GS21	`gvyzlg` – gvyzlf at γ_x 0.1	0.789	27%	0.718	+0.064	9	0.31	0.725	0.724	0.657	0.729	0.217	-0.0051 (t -1.7)	-0.06	3/10
GS21	`gvyp10c` – the product with a leverage-level term, exposure 0.10 $\tilde{b}(\tilde{z}+2)$	0.774	-86%	0.763	+0.006	1	0.09	0.745	0.727	0.720	0.767	0.674	-0.0221 (t -4.1)	-0.17	0/10
GS21	`gvytl` – three exposure types × three fixed leverage types (KP14's recipe)	0.739	84%	0.728	+0.009	1	0.12	0.708	0.688	0.685	0.714	0.621	-0.0060 (t -3.7)	-0.07	2/10
GS21	`gvytlg` – the same at γ_x 0.1	0.517	72%	0.498	+0.017	3	0.13	0.473	0.448	0.443	0.482	0.374	-0.0089 (t -4.7)	-0.09	0/10
BGN	`bvy9t2` – two priced chains, nine types (independent ladders), total price 2.5	0.636	42%	0.398	+0.063	16	0.23	0.395	0.339	0.355	0.383	0.266	+0.0124 (t 3.5)	+0.08	8/10
BGN	`bvychs` – bvych0 with sign-switching types: six types, loadings $\pm0.01, \pm0.035, \pm0.07$	0.510	36%	0.349	+0.069	20	0.23	0.360	0.274	0.277	0.334	0.183	+0.0255 (t 4.9)	+0.12	9/10
KP14	`vyg25s` – vyg25 with sign-switching types: six types, loadings $\pm0.02, \pm0.07, \pm0.14$	0.575	11%	0.174	+0.349	226	0.49	0.373	0.102	0.190	0.153	0.056	+0.1810 (t 17.2)	+0.32	10/10
KP14	`vyg25u` – vyg25 without the level compensation: the type shows in value at $y=0$	0.495	11%	0.351	+0.066	19	0.22	0.339	0.290	0.293	0.341	0.044	-0.0018 (t -0.5)	-0.02	3/10
GS21	`gvy25s` – gvy25 with sign-switching types: six level-exposure types, $\pm0.015, \pm0.05, \pm0.1$	0.619	52%	0.554	+0.061	11	0.27	0.549	0.541	0.506	0.560	0.318	-0.0118 (t -4.4)	-0.10	2/10
KP14	`kpsgn` – the price of x switches sign with y: $\gamma_x(2\,\mathrm{logistic}(-2y)-1)$ (conditioning room)	0.143	42%	0.075	+0.050	90	0.10	0.060	0.018	0.045	0.064	-0.002	-0.0041 (t -0.9)	-0.01	1/3
KP14	`vym6` – six priced states, 24 rank-6 types (vym3's construction at K = 6)	0.421	15%	0.300	+0.056	19	0.19	0.287	0.260	0.275	0.286	0.059	+0.0002 (t 0.1)	-0.01	4/10
KP14	`vym3s` – vyg25s's signed ladder over three priced states: 24 mean-zero types, mirrored alignments	0.503	13%	0.191	+0.224	126	0.37	0.296	0.134	0.195	0.171	0.065	+0.0998 (t 20.9)	+0.22	10/10
KP14	`vym6s` – vyg25s's signed ladder over six priced states: 24 mean-zero types, mirrored alignments, orthogonal off-γ directions	0.507	13%	0.192	+0.225	124	0.37	0.296	0.132	0.194	0.168	0.065	+0.0987 (t 15.0)	+0.22	10/10
GS21	`gvy2s` – gvy25s's signed ladder over two priced chains, cash-flow exposure only: twelve mean-zero types	0.625	53%	0.567	+0.046	8	0.23	0.540	0.542	0.505	0.564	0.333	-0.0246 (t -5.1)	-0.16	0/2
GS21	`gvytls` – gvytl with the exposure ladder mirrored: six signed exposure types × three fixed coupons, eighteen mean-zero types	0.627	51%	0.580	+0.044	8	0.23	0.580	0.493	0.544	0.588	0.313	-0.0074 (t -1.0)	-0.06	1/3
BGN	`bvychs4` – bvychs at bv_comp 0.4	0.612	29%	0.519	+0.059	12	0.25	0.497	0.468	0.465	0.497	0.177	-0.0035 (t -0.3)	-0.01	2/5
BGN	`bvychs6` – bvychs at bv_comp 0.6	0.656	28%	0.578	+0.047	8	0.23	0.547	0.493	0.512	0.545	0.183	+0.0013 (t 0.3)	+0.00	1/3
BGN	`bvychs8` – bvychs at bv_comp 0.8	0.684	28%	0.590	+0.044	7	0.23	0.577	0.501	0.468	0.568	0.189	+0.0086 (t 1.0)	+0.06	2/3
KP14	`vyg25s6` – vyg25s at bv_comp 0.6	0.576	11%	0.212	+0.305	147	0.47	0.337	0.183	0.131	0.185	0.062	+0.1424 (t 24.2)	+0.27	5/5
KP14	`vyg25s18` – vyg25s at bv_comp 1.8	0.574	11%	0.209	+0.312	167	0.47	0.387	0.106	0.172	0.193	0.062	+0.1854 (t 8.7)	+0.33	5/5
GS21	`gvyib2` – gvyib at gs_bib −0.2 (the −1.0 and −0.5 solves diverged: the firm picks the coupon that makes investment near-free)	0.327	96%	0.326	+0.001	0	0.03	0.319	0.298	0.293	0.322	0.314	-0.0027 (t -11.0)	-0.04	0/10
GS21	`gxq35` – gsbase with the x loading hump-shaped in productivity: $1-0.35\,(\tilde{z}^2-1)/\sqrt{2}$	0.284	96%	0.282	+0.002	1	0.03	0.270	0.248	0.242	0.270	0.269	+0.0002 (t 0.4)	+0.00	1/3
GS21	`gxq70` – gsbase with the x loading hump-shaped in productivity: $1-0.70\,(\tilde{z}^2-1)/\sqrt{2}$	0.278	95%	0.269	+0.008	3	0.07	0.261	0.211	0.201	0.261	0.259	+0.0006 (t 1.8)	+0.01	2/3
GS21	`gvy25sq` – gvy25s at γ_v 0.25, bv_comp 12.0: the level shifts kept, the y price cut	0.307	87%	0.305	+0.002	1	0.03	0.290	0.262	0.239	0.294	0.263	-0.0038 (t -8.2)	-0.05	0/3
GS21	`gvy25sh` – gvy25s at γ_v 0.5, bv_comp 9.0: the level shifts kept, the y price cut	0.351	71%	0.341	+0.009	3	0.08	0.328	0.309	0.286	0.336	0.246	-0.0077 (t -2.5)	-0.07	0/3
BGN	`bvychs18` – bvychs at bv_comp 1.8	0.340	49%	0.207	+0.055	27	0.16	0.218	0.198	0.178	0.208	0.167	+0.0064 (t 1.4)	+0.03	6/8
BGN	`bvychs25` – bvychs at bv_comp 2.5	0.449	36%	0.250	+0.163	67	0.33	0.341	0.286	0.223	0.246	0.159	+0.0482 (t 3.7)	+0.15	9/10
BGN	`bvychs30` – bvychs at bv_comp 3.0	0.488	32%	0.277	+0.182	67	0.36	0.384	0.318	0.268	0.279	0.157	+0.0626 (t 5.6)	+0.20	10/10
BGN	`bvychs35` – bvychs at bv_comp 3.5	0.508	31%	0.294	+0.189	67	0.38	0.403	0.315	0.309	0.303	0.156	+0.0712 (t 5.8)	+0.22	10/10
BGN	`bvychs40` – bvychs at bv_comp 4.0	0.518	30%	0.305	+0.189	65	0.39	0.413	0.305	0.334	0.319	0.155	+0.0590 (t 4.6)	+0.20	10/10
BGN	`bvychs45` – bvychs at bv_comp 4.5	0.525	29%	0.281	+0.224	81	0.42	0.422	0.316	0.345	0.315	0.154	+0.0638 (t 2.5)	+0.21	4/4
BGN	`bvychsn5` – bvychs at bv_comp -0.5	0.367	35%	0.254	+0.080	35	0.21	0.229	0.202	0.207	0.228	0.127	-0.0054 (t -0.9)	-0.01	6/10
BGN	`bvyvee` – V-shaped loading across a visible level ladder: level 0.6..1.4, b = +0.07, +0.00, -0.07, +0.00, +0.07	0.485	34%	0.352	+0.076	23	0.24	0.352	0.305	0.320	0.344	0.164	+0.0071 (t 1.3)	+0.05	6/8
GS21	`gvyql10` – U-shaped exposure in the firm's own leverage: $+0.1\,(\tilde{b}^2-1)$, one type	0.287	-81%	0.259	+0.015	8	0.08	0.247	0.230	0.231	0.249	0.238	-0.0021 (t -1.9)	-0.03	0/3
GS21	`gvyql20` – U-shaped exposure in the firm's own leverage: $+0.2\,(\tilde{b}^2-1)$, one type	0.602	-95%	0.593	+0.007	1	0.09	0.581	0.570	0.565	0.583	0.571	-0.0020 (t -1.9)	-0.05	0/3
GS21	`gvyqh10` – hump-shaped exposure in the firm's own leverage: $-0.1\,(\tilde{b}^2-1)$, one type	0.469	-66%	0.409	+0.032	11	0.16	0.366	0.358	0.335	0.391	0.339	-0.0259 (t -6.9)	-0.14	0/3
KP14	`vyg25sx2` – vyg25s with the loadings doubled: $\pm0.04, \pm0.14, \pm0.28$	0.640	6%	0.158	+0.437	343	0.57	0.456	0.206	0.105	0.134	0.027	+0.2452 (t 11.1)	+0.40	10/10
KP14	`vyg25sx3` – vyg25s with the loadings tripled: $\pm0.06, \pm0.21, \pm0.42$	0.661	5%	0.248	+0.364	175	0.55	0.461	0.223	0.098	0.237	0.021	+0.2177 (t 9.7)	+0.39	10/10
GS21	`gvyzz10` – U-shaped exposure in the firm's own productivity: $+0.1\,(\tilde{z}^2-1)$, one type	0.621	43%	0.356	+0.224	60	0.45	0.498	0.455	0.312	0.489	0.252	+0.0086 (t 0.7)	+0.02	3/5
GS21	`gvyzz20` – U-shaped exposure in the firm's own productivity: $+0.2\,(\tilde{z}^2-1)$, one type	0.686	26%	0.441	+0.168	43	0.42	0.521	0.350	0.321	0.462	0.179	+0.0589 (t 6.3)	+0.23	10/10
GS21	`gvyzzh10` – hump-shaped exposure in the firm's own productivity: $-0.1\,(\tilde{z}^2-1)$, one type	0.639	25%	0.364	+0.250	70	0.49	0.480	0.325	0.255	0.400	0.152	+0.0795 (t 4.6)	+0.25	10/10
GS21	`gvy25sc25` – gvy25s at bv_comp 2.5: strong level compensation	0.712	33%	0.542	+0.168	31	0.46	0.635	0.585	0.630	0.655	0.231	-0.0198 (t -2.8)	-0.14	1/6
BGN	`bvyhx2` – hump-shaped loading across a mild level ladder 0.85..1.15: b = -0.14, +0.00, +0.14, +0.00, -0.14	0.597	27%	0.526	+0.057	11	0.25	0.523	0.392	0.315	0.508	0.158	+0.0150 (t 0.7)	+0.06	2/3
BGN	`bvyhx3` – hump-shaped loading across a mild level ladder 0.8..1.2: b = -0.21, +0.00, +0.21, +0.00, -0.21	0.551	28%	0.349	+0.155	44	0.36	0.375	0.295	0.280	0.340	0.153	+0.0345 (t 3.9)	+0.14	10/10
BGN	`bvyhx4` – hump-shaped loading across a mild level ladder 0.8..1.2: b = -0.28, +0.00, +0.28, +0.00, -0.28	0.642	25%	0.357	+0.260	73	0.50	0.476	0.341	0.293	0.344	0.159	+0.1186 (t 10.9)	+0.31	10/10
KP14	`vyhx3` – hump-shaped loading across a mild level ladder 0.8..1.2: b = -0.21, +0.00, +0.21, +0.00, -0.21	0.632	7%	0.207	+0.377	197	0.54	0.431	0.217	0.136	0.178	0.031	+0.2123 (t 12.5)	+0.37	10/10
KP14	`vyhx4` – hump-shaped loading across a mild level ladder 0.8..1.2: b = -0.28, +0.00, +0.28, +0.00, -0.28	0.659	5%	0.298	+0.318	123	0.53	0.480	0.309	0.199	0.278	0.022	+0.1664 (t 10.6)	+0.36	10/10
GS21	`gvyhx3` – hump-shaped loading across a mild level ladder 0.8..1.2 (gs_level): b = -0.21, +0.00, +0.21, +0.00, -0.21	0.823	26%	0.706	+0.112	17	0.41	0.766	0.737	0.725	0.764	0.210	+0.0023 (t 0.9)	+0.04	7/10
BGN	`bvyq2` – two chains, the second loading quadratic in the first: $b_2=-14\,(b_1^2-\bar{b}_1^2)$, $b_1=\pm0.01, \pm0.035, \pm0.07$	0.460	37%	0.303	+0.063	21	0.20	0.307	0.240	0.260	0.285	0.172	+0.0223 (t 18.1)	+0.11	3/3
BGN	`bvyq2x` – two chains, the second loading quadratic in the first: $b_2=-7\,(b_1^2-\bar{b}_1^2)$, $b_1=\pm0.02, \pm0.07, \pm0.14$	0.523	30%	0.375	+0.070	19	0.24	0.359	0.321	0.308	0.366	0.158	-0.0075 (t -1.2)	-0.05	1/3
BGN	`bvyq2k` – two chains, the second loading quadratic in the first: $b_2=-28\,(b_1^2-\bar{b}_1^2)$, $b_1=\pm0.01, \pm0.035, \pm0.07$	0.470	34%	0.315	+0.071	22	0.22	0.304	0.265	0.261	0.302	0.159	+0.0026 (t 0.9)	+0.01	6/10
KP14	`vyg25su` – vyg25s without level compensation (bv_comp 0): $\pm0.02, \pm0.07, \pm0.14$	0.575	11%	0.290	+0.230	84	0.43	0.395	0.270	0.066	0.274	0.056	+0.1162 (t 11.1)	+0.28	10/10
KP14	`vyg25sux2` – vyg25sx2 without level compensation (bv_comp 0): $\pm0.04, \pm0.14, \pm0.28$	0.640	6%	0.185	+0.411	500	0.55	0.482	0.284	0.088	0.229	0.027	+0.1914 (t 6.5)	+0.37	10/10
KP14	`vyg25x2` – vyg25 (positive ladder) with the loadings x2: 0.04, 0.14, 0.28	0.468	7%	0.335	+0.074	22	0.23	–	–	–	–	–	–	–	–
KP14	`vyg25x3` – vyg25 (positive ladder) with the loadings x3: 0.06, 0.21, 0.42	0.475	6%	0.358	+0.063	17	0.22	–	–	–	–	–	–	–	–
KP14	`vyg35x2` – vyg25x2 (positive ladder x2) with the y price raised to 3.5	0.531	5%	0.373	+0.079	21	0.26	–	–	–	–	–	–	–	–
KP14	`vyg25x2c` – vyg25x2 (positive ladder x2) with the level compensation strengthened to 2.5	0.489	7%	0.359	+0.074	20	0.24	–	–	–	–	–	–	–	–
KP14	`vyge` – the y exposure on the profitability deviation, $b_f\,(\varepsilon-1)$: three types 0.1 / 0.25 / 0.45	0.477	19%	0.267	+0.038	14	0.15	0.250	0.228	0.254	0.251	0.088	-0.0048 (t -1.9)	-0.05	0/3
KP14	`vygel` – vyge with a visible level ladder, type_level 0.6 / 1.0 / 1.6	0.491	17%	0.269	+0.046	17	0.16	0.266	0.254	0.236	0.256	0.077	+0.0049 (t 1.1)	+0.03	2/3
KP14	`vygesm` – exposure on the profitability deviation, mirrored types ±0.05 / 0.1 / 0.15 with a level ladder 0.5..2.0: a product of two mean-zero factors	0.369	20%	0.232	+0.022	9	0.10	0.209	0.203	0.189	0.217	0.074	-0.0083 (t -1.4)	-0.04	1/3
KP14	`vygelb` – bounded exposure (cap 2) on the profitability deviation, mirrored types ±0.1 / 0.25 / 0.45 with a level ladder 0.5..2.0	0.482	16%	0.228	+0.058	26	0.17	0.219	0.199	0.203	0.215	0.074	+0.0024 (t 1.4)	+0.02	7/10
KP14	`vygelb7` – bounded exposure (cap 2), mirrored types ±0.2 / 0.45 / 0.7 with a level ladder 0.5..2.0	0.535	14%	0.231	+0.085	37	0.22	0.225	0.202	0.208	0.218	0.073	+0.0049 (t 2.1)	+0.03	8/10
KP14	`vygelb7g` – bounded exposure (cap 2), mirrored types ±0.2 / 0.45 / 0.7, level ladder 0.5..2.0, γ_v 4.0	0.677	11%	0.235	+0.120	51	0.27	0.238	0.208	0.208	0.222	0.073	+0.0135 (t 4.2)	+0.07	10/10
KP14	`vygelbw` – bounded exposure (cap 2), mirrored types ±0.3 / 0.7 / 1.2, level ladder 0.5..2.0, γ_v 4.0	0.710	11%	0.243	+0.138	56	0.29	0.254	0.216	0.216	0.228	0.073	+0.0215 (t 4.3)	+0.10	10/10
KP14	`vygelbx` – bounded exposure (cap 2), mirrored types ±0.5 / 1.2 / 2.0, level ladder 0.5..2.0, γ_v 4.0	0.727	11%	0.254	+0.150	58	0.31	0.273	0.227	0.216	0.241	0.073	+0.0297 (t 4.2)	+0.12	10/10
KP14	`vygelbc` – bounded exposure (cap 2.5), mirrored types ±0.3 / 0.7 / 1.2, level ladder 0.5..2.0, γ_v 4.0	0.752	9%	0.244	+0.170	69	0.33	0.272	0.220	0.219	0.231	0.068	+0.0369 (t 6.6)	+0.14	10/10
KP14	`vygelbc6` – bounded exposure (cap 2.5), mirrored types ±0.3 / 0.7 / 1.2, level ladder 0.5..2.0, γ_v 6.0	0.892	8%	0.249	+0.201	80	0.37	0.290	0.230	0.222	0.237	0.069	+0.0487 (t 5.8)	+0.16	10/10
KP14	`vygelbcs` – bounded exposure (cap 2.5), mirrored types ±0.3 / 0.7 / 1.2, level ladder 0.5..2.0, γ_v 4.0, σ_ϵ 0.3	0.858	7%	0.297	+0.256	90	0.46	0.388	0.251	0.288	0.282	0.052	+0.0913 (t 12.3)	+0.25	10/10
GS21	`gvylu30c`‡ – exposure U-shaped in leverage: fixed coupons 0.15 / 0.40 / 0.70, b = 0.3, 0, 0.3, value-compensated	1.219	99%	1.213	+0.005	0	0.11	–	–	–	–	–	–	–	–
GS21	`gvylu30`‡ – exposure U-shaped in leverage: fixed coupons 0.15 / 0.40 / 0.70, b = 0.3, 0, 0.3, no compensation	0.967	95%	0.951	+0.016	2	0.17	–	–	–	–	–	–	–	–
GS21	`gvyu4`‡ – U-shaped loading across a mild level ladder 0.8..1.2 (gs_level): b = +0.28, 0, 0, 0, +0.28	0.917	85%	0.893	+0.020	2	0.18	–	–	–	–	–	–	–	–
KP14	`vyu5r` – U-shaped loading across a mild level ladder 0.8..1.2: b = +0.35 at the ends, unexposed reference type	0.342	14%	0.267	+0.050	19	0.17	0.254	0.219	0.205	0.253	0.036	+0.0003 (t 0.2)	-0.00	4/10
KP14	`vyu4r` – U-shaped loading across a mild level ladder 0.8..1.2: b = +0.28 at the ends, unexposed reference type	0.367	14%	0.262	+0.069	27	0.20	0.251	0.223	0.208	0.248	0.038	+0.0025 (t 0.7)	+0.01	5/10
KP14	`vyu4` – U-shaped loading across a mild level ladder 0.8..1.2: b = +0.28, 0, 0, 0, +0.28	0.371	13%	0.263	+0.072	28	0.20	0.256	0.216	0.209	0.246	0.040	+0.0100 (t 1.5)	+0.04	7/10
BGN	`bvyu6r` – U-shaped loading across a mild level ladder 0.8..1.2: b = +0.42 at the ends, unexposed reference type	0.768	49%	0.592	+0.137	24	0.42	0.616	0.479	0.533	0.582	0.373	+0.0246 (t 2.6)	+0.13	8/10
BGN	`bvyu5r` – U-shaped loading across a mild level ladder 0.8..1.2: b = +0.35 at the ends, unexposed reference type	0.774	46%	0.641	+0.108	17	0.38	0.664	0.562	0.558	0.623	0.358	+0.0295 (t 3.1)	+0.16	9/10
BGN	`bvyu4` – U-shaped loading across a mild level ladder 0.8..1.2: b = +0.28, 0, 0, 0, +0.28	0.382	31%	0.246	+0.111	54	0.25	0.241	0.157	0.181	0.227	0.118	+0.0142 (t 1.3)	+0.06	7/10
KP14	`vyg25` – three types on y, price 2.5 (control)	0.495	11%	0.346	+0.082	24	0.25	0.350	0.292	0.320	0.335	0.044	+0.0148 (t 5.2)	+0.09	10/10
GS21	`gvyp10` – product exposure $\tilde{z}\tilde{b}$ at 0.10 (control)	0.650	8%	0.447	+0.194	47	0.46	0.553	0.238	0.333	0.548	0.018	+0.0057 (t 0.7)	+0.01	5/10
BGN	`bvych0` – y as a 21-node chain, three types (control)	0.576	37%	0.341	+0.055	16	0.20	0.334	0.285	0.292	0.326	0.216	+0.0076 (t 2.7)	+0.05	8/10
GS21	`gvysl15`‡ – industry sign × leverage: two industries (gs_level 0.8 / 1.2) × fixed coupons 0.15 / 0.40 / 0.70, b = sign × 0.15 × coupon/0.4	1.190	66%	1.183	+0.006	1	0.12	–	–	–	–	–	–	–	–
GS21	`gvysz2` – type sign × productivity $\tilde{z}$: mirrored types ±0.15 / 0.3 / 0.5, gs_level ladder 0.5..2.0, no compensation	0.847	17%	0.603	+0.208	37	0.54	0.678	0.536	0.488	0.619	0.151	+0.0590 (t 11.5)	+0.27	10/10
GS21	`gvysz1` – type sign × productivity $\tilde{z}$: mirrored types ±0.1 / 0.2 / 0.35, gs_level ladder 0.5..2.0, no compensation	0.844	25%	0.530	+0.278	60	0.60	0.658	0.560	0.431	0.605	0.210	+0.0533 (t 6.1)	+0.24	10/10
BGN	`bvychsx2` – bvychs with the loadings x 2: $\pm0.02, \pm0.07, \pm0.14$	0.496	32%	0.344	+0.089	26	0.26	0.347	0.308	0.279	0.333	0.161	+0.0072 (t 1.1)	+0.04	7/10
BGN	`bvychsx3` – bvychs with the loadings x 3: $\pm0.03, \pm0.105, \pm0.21$	0.533	31%	0.384	+0.081	21	0.26	0.377	0.353	0.321	0.370	0.164	+0.0024 (t 0.5)	+0.01	6/10
GS21	`gvy25sx2` – gvy25s with the loadings x 2: $\pm0.03, \pm0.1, \pm0.2$	0.817	41%	0.721	+0.093	13	0.38	0.814	0.777	0.796	0.814	0.344	+0.0005 (t 0.1)	+0.01	2/3
KP14	`kpgam` – $\gamma_x\times$ logistic(y), 1 to 3	0.390	65%	0.373	+0.005	1	0.06	0.358	0.348	0.340	0.363	0.254	-0.0053 (t -4.0)	-0.05	1/10
KP14	`vyx` – three types on y, price 1.8	0.421	14%	0.321	+0.055	17	0.19	0.313	0.284	0.302	0.310	0.054	+0.0009 (t 0.3)	+0.02	7/10
KP14	`vyj25` – the same with a priced crash in y	0.554	10%	0.323	+0.104	33	0.28	0.315	0.270	0.292	0.307	0.045	+0.0080 (t 2.2)	+0.06	8/10
KP14	`vym2t3` – two priced states, three loading vectors	0.475	12%	0.326	+0.090	28	0.26	0.317	0.280	0.301	0.308	0.049	+0.0075 (t 2.1)	+0.05	8/10
KP14	`vy9t2f` – two priced states, nine types, each state priced 2.5	0.582	6%	0.415	+0.077	19	0.26	0.416	0.314	0.330	0.401	0.018	+0.0149 (t 4.2)	+0.10	9/10
KP14	`vym3t3` – three priced states, three types	0.443	14%	0.285	+0.087	30	0.24	0.279	0.246	0.266	0.273	0.061	+0.0036 (t 0.8)	+0.02	6/10
KP14	`vym3` – three priced states, twenty rank-3 types	0.418	15%	0.298	+0.056	19	0.19	0.286	0.259	0.278	0.287	0.059	-0.0032 (t -1.7)	-0.02	5/10
RF	`kfh_tanh_1` – K=1, bounded (tanh) loadings on observed characteristics	0.592	8%	0.592	+0.001	0	0.03	0.579	0.578	0.570	0.588	0.041	-0.0105 (t -1.9)	-0.10	0/5
RF	`kfh_tanh_4` – K=4	0.584	9%	0.574	+0.007	1	0.09	0.541	0.564	0.544	0.560	0.041	-0.0232 (t -2.9)	-0.16	0/5
RF	`kfh_tanh_16` – K=16	0.469	11%	0.378	+0.044	12	0.19	0.346	0.357	0.345	0.359	0.041	-0.0152 (t -1.4)	-0.10	1/5
RF	`lsh_a0_b0_n0` – K=1 product of two mean-zero latent states (sign ladder origin), no noise	0.598	8%	0.204	+0.393	192	0.56	0.569	0.154	0.538	0.175	0.041	+0.0309 (t 3.6)	+0.18	5/5
RF	`lsh_a0_b1_n0` – one factor's mean shifted 1 sd	0.598	8%	0.594	+0.004	1	0.07	0.582	0.579	0.577	0.593	0.041	-0.0110 (t -1.8)	-0.11	1/5
RF	`lsh_a0_b2_n0` – one factor shifted 2 sd	0.598	8%	0.597	+0.001	0	0.03	0.587	0.582	0.578	0.596	0.041	-0.0094 (t -1.9)	-0.10	0/5
RF	`lsh_a1_b1_n0` – both factors shifted 1 sd	0.598	8%	0.596	+0.002	0	0.05	0.588	0.582	0.584	0.592	0.041	-0.0048 (t -2.0)	-0.07	0/5
RF	`lsh_a2_b2_n0` – both 2 sd	0.598	8%	0.597	+0.001	0	0.02	0.590	0.583	0.583	0.594	0.041	-0.0047 (t -1.9)	-0.07	0/5
RF	`lsh_a3_b3_n0` – both 3 sd	0.598	8%	0.598	+0.000	0	0.02	0.591	0.583	0.583	0.595	0.041	-0.0046 (t -1.9)	-0.07	0/5
RF	`lsh_a0_b0_n1` – mean-zero product, observation noise 1.0	0.598	8%	0.160	+0.422	263	0.56	0.524	0.094	0.431	0.133	0.041	+0.0928 (t 3.7)	+0.30	5/5
RF	`lsh_a2_b2_n1` – both shifted 2 sd, noise 1.0	0.598	8%	0.593	+0.001	0	0.03	0.583	0.582	0.572	0.589	0.041	-0.0060 (t -2.1)	-0.08	0/5
RF	`lrt_n0` – one-signed coverage ratio of the two latents, no noise	0.598	8%	0.597	+0.001	0	0.04	0.591	0.579	0.571	0.593	0.041	-0.0033 (t -2.3)	-0.05	0/5
RF	`lsq_a0_n0` – K=1 square of the SIZE latent, mean zero, no noise (the VW market holds size^2)	0.598	8%	0.175	+0.422	242	0.57	0.557	0.153	0.565	0.151	0.041	-0.0082 (t -1.5)	-0.10	2/5
RF	`lsq_a0_n1` – noise 1.0	0.598	8%	0.144	+0.436	302	0.56	0.508	0.119	0.540	0.130	0.041	-0.0313 (t -2.0)	-0.18	1/5
RF	`lsq_a1_n0` – size latent shifted 1 sd	0.598	8%	0.596	+0.002	0	0.05	0.585	0.578	0.581	0.595	0.041	-0.0101 (t -1.9)	-0.11	0/5
RF	`lsq_a2_n0` – shifted 2 sd	0.598	8%	0.598	+0.000	0	0.02	0.588	0.580	0.579	0.597	0.041	-0.0089 (t -1.9)	-0.10	0/5
RF	`lsqb_a1_n0` – bm latent shifted 1 sd	0.598	8%	0.596	+0.002	0	0.05	0.590	0.582	0.577	0.593	0.041	-0.0027 (t -2.2)	-0.06	0/5
RF	`lsqb_a1_n1` – shifted 1 sd, noise 1.0	0.598	8%	0.589	+0.002	0	0.05	0.581	0.579	0.560	0.585	0.041	-0.0047 (t -1.9)	-0.07	1/5
RF	`lsqb_a2_n0` – shifted 2 sd	0.598	8%	0.597	+0.000	0	0.02	0.592	0.583	0.580	0.594	0.041	-0.0026 (t -2.3)	-0.05	0/5
RF	`lsqb_a2_n1` – shifted 2 sd, noise 1.0	0.598	8%	0.593	+0.001	0	0.03	0.584	0.582	0.568	0.589	0.041	-0.0047 (t -1.9)	-0.07	0/5

- `bvyc0`: `bvy25` with $\theta_f = 1$, so the type is visible in value.

- `bvyj25`: `bvych0` (group 3b) plus a crash: each month, with probability 1% whatever y_t, y_{t+1} goes to the lowest node; otherwise the Tauchen transition. The kernel weights the crash transition by e^2, so its risk-neutral probability is about 7.4% a month before renormalization.

- `bvysg`: `bgnbase` plus y; type f's project cash flows are $C_j e^{b_f y_t}$, $b_f \in \{-0.07, 0, +0.07\}$, not level-compensated.

- `bvychs`: `bvych0` with the ladder mirrored: six types of equal share, $b_f \in \{-0.07, -0.035, -0.01, 0.01, 0.035, 0.07\}$, level-compensated (bv_comp 1.2, θ_f 1 to 1.60 with $b = -0.07$ the reference), 21-node chain, price 2.5.

- `bvychs18`, `bvychs25`, `bvychs30`, `bvychs35`, `bvychsn5`: `bvychs` at bv_comp 1.8 / 2.5 / 3.0 / 3.5 / −0.5. Ten seeds each (bvychs18 eight). bgn_vy's SDF step had its K4 pair sum rewritten as a blocked sum (commit 78961c4) after these ran out of memory at 36 GB; the arithmetic is unchanged to 3×10^{-15}.

- `bvyvee`: five types of equal share on a visible level ladder, $\theta_f = 0.6$ / 0.8 / 1.0 / 1.2 / 1.4 multiplying the cash-flow level, exposure V-shaped across it, $b_f = +0.07$, 0, −0.07, 0, +0.07; bv_comp 1.0, price 2.5. Eight seeds; two stopped. `bvyhump`, its mirror, was solved and not simulated.

- `bvychsx2`, `bvychsx3`: `bvychs` with the loadings ×2 (±0.02, ±0.07, ±0.14) and ×3 (±0.03, ±0.105, ±0.21). Ten seeds each.

- `bvychs40`, `bvychs45`: `bvychs` at bv_comp 4.0 / 4.5. Ten seeds; four seeds, stopped.

- `bvyhx2`, `bvyhx3`, `bvyhx4`: five types on a mild visible level ladder (θ_f 0.85–1.15 / 0.80–1.20 / 0.80–1.20) with the loading hump-shaped across it, $b_f = -c, 0, +c, 0, -c$, $c = 0.14$ / 0.21 / 0.28; bv_comp 1, price 2.5. Three seeds (stopped, room 9–13%); ten; ten.

- `bvyu4`, `bvyu5`, `bvyu6`, `bvyu5r`, `bvyu6r`: the hump ladder's chassis with a U-shaped loading, $+c$ at the two ends, $c = 0.28$ / 0.35 / 0.42; `bvyu5` failed the gate (the small exposed type is the compensation's reference, $J^* = 0$), `bvyu6` was not built, the `r` versions list the unexposed level-1.0 type first. Ten seeds each.

- `bvyq2`, `bvyq2x`, `bvyq2k`: `bvychs`'s six types with a second priced chain (both at $2.5/\sqrt{2}$) whose loading is quadratic in the first, $b_2 = \kappa\,(b_1^2 - \overline{b_1^2})$: $\kappa = -14$ on ±0.01/0.035/0.07; −7 on ±0.02/0.07/0.14; −28 on ±0.01/0.035/0.07. Three seeds each (stopped); ten.

GS21

- `g28`: `gsbase` with $\gamma_x \to \gamma(x_t) = \text{clip}\big(0.5 - 0.28\,x_t/\text{sd}(x),\ 0.05,\ 1.0\big)$, countercyclical.

- `bx7`: `gsbase` with $\gamma_x \to g_s\gamma_x$, $g \in \{0.6, 3.0\}$ on a two-state Markov regime, and five types with output $e^{b_f x_t + z_{it}} k_{it}$, $b_f \in \{1, 2.5, 4, 5.5, 7\}$.

- `gvyp10`: `gsbase` plus y (9-node chain), one type, output $e^{x_t + z_{it} + 0.10\,\tilde{z}_{it}\tilde{b}_{it}\,y_t} k_{it}$, with $\tilde{z} = z/\text{sd}(z)$ and $\tilde{b} = (b - 0.5)/0.29$ the standardized coupon. Exposure is productivity × leverage, revealed by profitability and leverage.

- `gx7`: `gsbase` with five types, output $e^{b_f x_t+z_{it}}k_{it}$, $b_f \in \{1, 2.5, 4, 5.5, 7\}$, and $\gamma(x_t) = \mathrm{clip}\big(0.5 - 0.4\,x_t/\mathrm{sd}(x),\ 0.05,\ 1.5\big)$. A firm's exposure to the one priced shock is its type, which size and profitability reveal.

- `gvy1`: `gsbase` plus y (9-node chain over ±2.5 sd), one type; output $e^{x_t+z_{it}+0.05\,y_t}k_{it}$, a level exposure.

- `gvy1j`: `gvy1` with $\gamma_v = 0$ and the crash, on a chain over ±3.5 sd.

- `gvy25`: `gsbase` plus y (9-node chain); output $e^{x_t+z_{it}+b_f y_t+a_f}k_{it}$, $b_f \in \{0.015, 0.05, 0.10\}$, with a_f the level compensation.

- `gvy25s`: `gvy25` with the ladder mirrored: six types of equal share, $b_f \in \{-0.10, -0.05, -0.015, 0.015, 0.05, 0.10\}$, level-compensated (bv_comp 1.2) at price 2.5.

- `gvys25`: `gsbase` plus y, entering through an operating cost: $0.2\,e^{c_f y_t}$ per unit capital is subtracted from output, $c_f \in \{+0.4, 0, -0.4\}$; output itself does not load on y.

- `gvyp05, gvyp075, gvyp15, gvyp20`: `gvyp10` with the coefficient 0.10 replaced by 0.05, 0.075, 0.15 and 0.20.

- `gvyp10x`: `gvyp10` with output $e^{(1+(\tilde{z}_{it}^2-1)/\sqrt{2})\,x_t+z_{it}+0.10\,\tilde{z}_{it}\tilde{b}_{it}\,y_t}k_{it}$: the loading on x is also nonlinear in productivity.

- `gvyp10j`: `gvyp10` (group 3b) plus the same priced crash on its 9-node chain.

- `gvypk2`: `gvyp10` with a second independent chain (5 nodes each, 25 joint); two types of equal share with output $e^{x_t+z_{it}+0.10\,\tilde{z}_{it}\tilde{b}_{it}\,y_t}k_{it}$ and $e^{x_t+z_{it}+0.10\,\tilde{z}_{it}\tilde{b}_{it}\,y_t^{(2)}}k_{it}$.

- `gvyzl`: `gvyz10` with three leverage types: the coupon the firm refinances to is fixed at 0.158 / 0.421 / 0.684 per unit capital, equal shares drawn independently of productivity; refinancing opportunities (1% a month) and dilution at investment as in GS21, so the realized coupon drifts below the type's.

- `gvyzlf`: `gvyzl` with the capital structure fixed for life: no refinancing, new debt issued pari passu at investment (debt-to-equity 0.25 / 1.0 / 3.25).

- `gvyzlw`: `gvyzlf` with coupons 0.05 / 0.45 / 0.85; the high type is distressed (premium 16% a year, default rate 1.4%).

- `gvyzl5`: `gvyzlf` with five types, coupons 0.10 to 0.70.

- `gvyzlb`: `gvyzlf` with the exposure doubled to $0.20\,\tilde{z}\,y$.

- `gvyzlr`: `gvyzlf` with riskless debt: the debt claim is worth its riskless coupon value in every state (creditors made whole on default).

- `gvyzlp`: `gvyzlf` with $\kappa_y = 0.10$ a year instead of 0.35.

- `gvyzlrp`: both.

- `gsbg1`: the GS21 benchmark with $\gamma_x = 0.1$ instead of 0.5 (`gs_vy` at one y node and no exposure, which is `gsbase`'s solver).

- `gvyz10g`: `gvyz10` at $\gamma_x = 0.1$.

- `gvyzlg`: `gvyzlf` at $\gamma_x = 0.1$.
- `gvyp10c`: `gvyp10` with a leverage-level term, output $e^{x+z+0.10\,\tilde{b}(\tilde{z}+2)\,y}$: a levered firm's exposure keeps leverage's sign for every firm with $\tilde{z} > -2$. The mean firm ($\tilde{b} < 0$) is a hedge against y, so the market's premium is negative.
- `gvytl`: KP14's 3b recipe in GS21: three latent exposure types, output $e^{x+z+b_f y+a_f}$ with $b_f \in \{0.015, 0.05, 0.10\}$ and a_f the level compensation, times three fixed leverage types (coupons 0.158 / 0.421 / 0.684, no refinancing), nine types of equal share.
- `gvytlg`: `gvytl` at $\gamma_x = 0.1$ (benchmark `gsbg1`).
- `gxq35`, `gxq70`: `gsbase` (no y) with the productivity loading $1 + c\,(\tilde{z}^2 - 1)/\sqrt{2}$, $c = -0.35$ / -0.70. Ten oracle seeds, three estimator seeds; stopped.
- `gvy25sq`, `gvy25sh`: `gvy25s` at γ_v 0.25 / 0.5 with bv_comp 12 / 9 (the types' level shifts kept, the y price cut). Ten oracle seeds, three estimator seeds; stopped.
- `gvyql10`, `gvyql20`, `gvyqh10`: one type on `gvy25s`'s chassis with the y exposure $c\,(\tilde{b}^2 - 1)$, $\tilde{b}$ the standardized coupon, $c = +0.1$ / $+0.2$ / -0.1. Three seeds each; stopped.
- `gvy25sx2`: `gvy25s` with the loadings ×2 (± 0.03, ± 0.10, ± 0.20). Five oracle seeds, three estimator seeds; stopped. `gvy25sx3`: the $b_f = +0.045$ type's solve oscillated to the sweep cap; not run.
- `gvyzz10`, `gvyzz20`, `gvyzzh10`: one type on `gvy25s`'s chassis with the y exposure $c\,(\tilde{z}^2 - 1)$, $\tilde{z}$ the standardized idiosyncratic productivity, $c = +0.1$ / $+0.2$ / -0.1 (`gs_bzz`, new). Five (stopped), ten and ten seeds.
- `gvyhx3`, `gvyhx4`, `gvyu4`: the size ladder in GS21 through the new gs_level (a per-type production-level multiplier; 1.0 leaves every earlier economy unchanged): five types at 0.8–1.2, bv_comp 1.0, NY 9, y_mult 2.5, price 2.5, one solve per type; the hump ±0.21 / ±0.28 and the U +0.28. gvyhx3 ten seeds (no gap), gvyhx4 solved and not pursued, gvyu4 stopped at the oracle.
- `gvylu30`, `gvylu30c`: the exposure U-shaped in leverage on `gvyzlf`'s fixed-coupon chassis (coupons 0.15 / 0.40 / 0.70 for life), production exposure gs_bv (0.3, 0, 0.3), without and with the $y = 0$ value compensation; both stopped at the oracle.
- `gvysz1`, `gvysz2`: the industry sign times the firm's standardized productivity (gs_bzl 0, gs_bzi 1), six mirrored types ±0.1/0.2/0.35 and ±0.15/0.3/0.5 on the gs_level ladder 0.5–2.0, bv_comp 0, price 2.5, six solves each; ten seeds each.
- `gvysl15`, `gvysl20`: the industry sign times financial leverage, two industries (gs_level 0.8 / 1.2) times three fixed coupons, exposure sign $\times c \times$ coupon/0.40 at $c = 0.15$ / 0.20, bv_comp 1.0; gvysl15 stopped at the oracle, gvysl20 solved and not run.
- `gvy25sc25`: `gvy25s` at bv_comp 2.5. Six seeds; stopped. `gvy25sc32`: one type's solve cycle-capped; not run.

KP14

- `kpgam`: `kpbase` with $\gamma_x \to \gamma_x\, g(y_t)$, $g(y) = 1 + 2/(1 + e^{2y})$, rising from 1 when y is high to 3 when it is low; y is unpriced ($\gamma_v = 0$) and no firm is exposed to it.

- `vyg25`: `vy1` with three types, $b_f \in \{0.02, 0.07, 0.14\}$, level-compensated. Growth options are options on the $e^{b_f y}$ cash flows, so a firm's exposure is b_f times the elasticity of its value, which rises with its growth-option share: type × growth-option share, revealed by size and book-to-market.
- `vyx`: `kpbase` plus y with $\gamma_v = 1.8$; type f's project cash flows are $\theta_f e^{b_f y_t} x_t u_{it} \varepsilon_{jt} K_j^{\alpha}$, $b_f \in \{0.02, 0.07, 0.14\}$, level-compensated.
- `vyj25`: `vyg25` (group 3b) plus a Poisson jump of y to $-y_{\max} = -3.5$ at intensity 0.12 a year (1% a month) under P and $0.12\, e^2 = 0.89$ a year under Q.
- `vym3t3`: `vyg25` with $y \in \mathbb{R}^3$, three independent OU states; type f's cash flows load $e^{b_f' y_t}$ with three vectors $b_f \in \mathbb{R}^3$ spanning the three states.
- `vym2t3`: `vyg25` with $y \in \mathbb{R}^2$, two independent OU states; three types with loading vectors $(0.02, 0)$, $(0, 0.07)$, $(0.099, 0.099)$, price $2.5/\sqrt{2}$ on each: three sorts in three directions, which prices types 1 and 2 at an effective 1.77 instead of 2.5. The group-4 cell until October 21, 2026.
- `vy9t2f`: `vy9t2` with each state priced at 2.5 instead of $2.5/\sqrt{2}$; nests `vyg25` exactly.
- `vyg25s`: `vyg25` with the ladder mirrored: six types of equal share, $b_f \in \{-0.14, -0.07, -0.02, 0.02, 0.07, 0.14\}$, level-compensated as vyg25 (bv_comp 1.2, θ_f 1 to 3.36 with $b = -0.14$ the reference), price 2.5.
- `vyg25u`: `vyg25` with bv_comp 0: no level compensation, so the type shows in value at $y = 0$. The returns are vyg25's exactly ($\mathrm{SR}_{\max}$ equal to every digit); only the characteristics change.
- `vym3`: the same three states with twenty types, $b_f \in \mathbb{R}^3$ of rank 3.

What they add to the main results:

- **The chain's resolution is not what separates `bvych0` from `bvy25`**: on 41 nodes the gap is +0.009 against +0.008 on 21, and the difference from the Gaussian economy's +0.015 (sd 0.016) is within noise.
- **A level exposure of GS21 output to y** (`gvy1`, `gvy1j`) has no room at all, as every other level wiring in GS21; the productivity exposure (`gvyz10`) has 5% and the product with the coupon 43%.
- **The crash on an interaction economy** leaves room and gap at the control's in BGN (+0.007 against +0.008) and GS21 (+0.030 against +0.020, within noise) and halves KP14's (+0.008 against +0.014): the crash months are a state the ranks-times-state linear benchmark reads as well as DKKM does.
- **The GS21 product ladder**: room is hump-shaped in the product's size (45 / 111 / 47 / 22 / 6% of the linear ceiling at 0.05 / 0.075 / 0.10 / 0.15 / 0.20) while the gap is +0.019 / +0.077 / +0.020 / +0.078 / +0.022; above 0.10 the market becomes a hedge. A second nonlinear loading on x (`gvyp10x`) gives +0.049: a second interaction factor adds. `bvy25`, the Gaussian-y version of `bvych0`, has twice its gap at the same room, and the visibility ladder (`bvy25`, `bvyc0`, and `bvyc10` at `bv_comp` 1.0: +0.015 / +0.010 / +0.010) is flat within noise.
- **Two apparent failures of the mechanism are estimation cost, not structure.** `vyx` is `vyg25` at a lower price of y: its room (0.054) is about what DKKM gives up to estimation relative to the linear model (0.053), so the gap is zero, and the higher price lifts the room more than the cost. `bvysg` has twice `bvy25`'s room but DKKM's estimation loss is two and a half times

larger (0.16 against 0.06): with signed exposures the nonlinear part of the tangency portfolio is a direct bet on the persistent state y, whose sample mean a 360-month window barely pins down. `bvychs` (appendix), the signed ladder with six types and the level compensation, has a positive gap of +0.026 at room 20%, so the loss in `bvysg` came with the visible type, as in `vyg25u`. What matters is room net of DKKM's estimation cost. Four more GS21 wirings (investment cost, floating coupon, a larger option, state-dependent recovery) were stopped at the oracle stage with room below +0.01, and GS21's default channel is dormant at every calibration tried.

- **Knobs that are not mechanisms**: the speed of the priced state (κ 0.15 / 0.35 / 0.70 in KP14: +0.007 / +0.014 / +0.009, a mild hump within noise), the visibility of the type, the number of types (`vym3` against `vym3t3`), the sample size (`vym3` at $N = 1000$: room up, gap flat), and the ridge grid (DKKM +0.0001). The price of the interacted component is the one scaling lever (`vyx` to `vyg25`).
- `vyg25sx2`, `vyg25sx3`: `vyg25s` with the loadings ×2 (±0.04, ±0.14, ±0.28; θ_f 1 to 11.2) and ×3 (±0.06, ±0.21, ±0.42; θ_f 1 to 37). Ten seeds each.
- `vyg25su`, `vyg25sux2`: `vyg25s` and `vyg25sx2` at bv_comp 0 (the parents' returns; the type revealed by value). Ten seeds each.
- `vyge`, `vygel`, `vygels`, `vygesm`: the y exposure on the profitability deviation (`bv_on_eps`, cash flow $xK^{\alpha}\theta u[1 + (\varepsilon - 1)e^{b_f y}]$): three types 0.1 / 0.25 / 0.45 with $\theta_f = 1$; the same with `type_level` 0.6 / 1.0 / 1.6; six mirrored types ±0.1/0.25/0.45 on a level ladder 0.5–2.0 (solved, not run: negative unit values for $b_f \leq -0.25$); ±0.05/0.1/0.15 on the same ladder. Three seeds each; stopped at the oracle.
- `vygelb`, `vygelb7`, `vygelb7g`, `vygelbw`, `vygelbx`, `vygelbc`, `vygelbc6`, `vygelbcs`: the bounded exposure on the profitability deviation (`bv_eps_cap` c: $h_b(y) = c/(1 + (c - 1)e^{-cby/(c-1)})$ in place of $e^{b_f y}$), six mirrored types on the level ladder 0.5–2.0, bv_comp 1.2: ±0.1/0.25/0.45 (cap 2, γ_v 2.5); ±0.2/0.45/0.7 (γ_v 2.5 and 4.0); ±0.3/0.7/1.2 (γ_v 4.0; cap 2 and 2.5); ±0.5/1.2/2.0 (γ_v 4.0); cap 2.5 at γ_v 6.0 or σ_ε 0.3. Ten seeds each.
- `vyhx3`, `vyhx4`, `vyu4`, `vyu4r`, `vyu5r`: BGN's size ladder in KP14 through type_level: five types at 0.8–1.2, bv_comp 1.0, price 2.5; the hump ±0.21 / ±0.28 and the U +0.28 (vyu4r / vyu5r with the unexposed level-1.0 type first, +0.28 / +0.35). Ten seeds each.
- `vyg25x2`, `vyg25x3`, `vyg35x2`, `vyg25x2c`: `vyg25` with the loadings ×2 / ×3, and ×2 at γ_v 3.5 / bv_comp 2.5. Three to six oracle seeds; stopped.

Two priced states

A second independent priced state, with the total price held at the one-state value ($2.5/\sqrt{2}$ on each). In BGN and GS21 three types span the two states with loading vectors in the pattern $(1, 0)$, $(0, 1)$, $(1, 1)/\sqrt{2}$, at the 3b exposure sizes; in GS21 each type carries the 0.10 product exposure to its state. In KP14 the cell is the nested extension of 3b: every firm draws a type for each state independently from the 3b ladder $(0.02, 0.07, 0.14)$, nine types in all, so exposure to each state is its own type times leverage and zeroing the second ladder recovers `vyg25`. The three-direction KP14 version (`vym2t3`), the nested version with each state priced at the full 2.5 (`vy9t2f`) and the nested BGN version (`bvy9t2`) are in the appendix.

What changes. Total price fixed, $\gamma_v = 2.5/\sqrt{2}$ on each.

Table 9: Group 4. Larger factor models: two priced states. Two independent priced states at total price 2.5 ($2.5/\sqrt{2}$ each). BGN and GS21: three types with loading vectors (1,0), (0,1), $(1,1)/\sqrt{2}$. KP14: nine types, an independent draw from the 3b ladder for each state (the three-direction version and the full-price version are in the appendix).

model	economy	$\mathrm{SR}_{\max}$	market share	linear ceiling	room	room % of linear	interaction SR	DKKM	FMR	FF	linear	CAPM	fair gap	DKKM ⊥	+
BGN	`bvyk2` – two priced chains, three loading vectors	0.639	36%	0.393	+0.071	18	0.25	0.394	0.332	0.353	0.379	0.228	+0.0152 (t 4.3)	+0.10	10/10
GS21	`gvypk3` – two priced chains, three product types	0.612	8%	0.349	+0.226	65	0.45	0.452	0.242	0.300	0.398	0.035	+0.0515 (t 3.4)	+0.18	9/10
KP14	`vy9t2` – two priced states, nine types (independent ladders), total price 2.5	0.506	9%	0.387	+0.056	15	0.21	0.378	0.326	0.337	0.373	0.031	+0.0048 (t 2.3)	+0.04	7/10

- `bvyk2`: `bvych0` with a second independent 21-node chain $y^{(2)}$; project cash flows

$$C_j\, e^{b_f y_t + b_f^{(2)} y_t^{(2)}}, \qquad (b_f, b_f^{(2)}) \in \{(0.01, 0),\, (0, 0.035),\, (0.049, 0.049)\}.$$

- `gvypk3`: `gvyp10` with a second independent chain (5 nodes each, 25 joint) and three types of equal share, loading $0.10\,\tilde{z}\tilde{b}$ on y, on $y^{(2)}$, and $0.071\,\tilde{z}\tilde{b}$ on both.
- `vy9t2`: `vyg25` with $y \in \mathbb{R}^2$, two independent OU states; a firm's type is a pair (f_1, f_2) drawn independently from $\{0.02, 0.07, 0.14\}$ for each state, nine types of equal share, and its cash flows load $e^{b_{f_1} y_t + b_{f_2} y_t^{(2)}}$; price $2.5/\sqrt{2}$ on each state.

Verdict. More factors are a lever only where the second state brings a second interaction. In BGN the second chain raises room by a quarter and doubles the gap (+0.015, ten of ten seeds). In GS21 it raises the product economy's gap from +0.020 to +0.056 (nineteen of twenty seeds), as the two-type version did (appendix). In KP14 it does not help however it is wired: the nested nine-type cell lowers room from 24% to 15% of the linear ceiling and the gap from +0.015 to +0.005; the three-direction version (`vym2t3`) keeps the room (28%) but halves the gap (+0.008); pricing each state at the full 2.5 (`vy9t2f`) only returns the gap to 3b's (+0.015, nine of ten seeds) at lower room (19%). The reason is in the ceilings: with an independent second type a firm's total exposure is the sum of two ladder draws, six evenly spread values that size and momentum proxy almost linearly, so the linear ceiling rises more than the nonlinear one (0.346 to 0.387). Both states act through the same channel, the growth-option elasticity, so the second state adds premium and a factor but not a second interaction. BGN's nine-type version (`bvy9t2`, +0.012, room 16%) sits between 3b and the three-direction cell. The reduced form (Table 7) says the same with the structure controlled.

Reduced form: bounded loadings, the sign ladder and the shifted squares

The RF rows of Table 8 are the `kfh` ladder with bounded (tanh) loadings, the sign ladder and the shifted squares of a latent; the curvature ladders are in Table 7.

The sign ladder (`lsh` rows) shifts the mean of one or both of the two latent factors in the $K = 1$ product by 0, 1, 2 or 3 standard deviations, and `lrt` replaces the product by a one-signed ratio of the two (a debt-service coverage ratio with a ridge). The mean-zero product has room of +0.39 (no noise) and +0.42 (noise 1.0) and fair gaps of +0.031 and +0.093, the $K = 1$ latent-product cell of the reduced-form table; every shifted cell and the ratio have room of 0.000 to 0.005 and a fair gap of −0.005 to −0.014, at the same $\mathrm{SR}_{\max}$ of 0.598. The linear-in-latents R^2 of the loading goes from 0.01 at no shift to 0.79 at a two-sd shift of one factor and 0.89 at two-sd shifts of both. Room in a one-factor economy is therefore all or nothing: a linear portfolio with any nonzero covariance with the exposure mimics the factor, diversifies the rest and takes its whole Sharpe, and only an exposure orthogonal to every linear function of the characteristics leaves the factor to the nonlinear estimator. This is what the structural sign-switching economies of Section 3b test.